

Local Resolution and Abstention:

A Controlled Comparison of Anisotropic Resolution Fields with Local-Metric and Prototype Baselines under a Predefined Criterion

Luciano Silva Alarco

Pontificia Universidad Católica del Perú

ORCID 0000-0002-3395-3129

Working draft v0.2 (after one round of internal review) — research line HQF001 — 3 October 2026

Abstract

Research line HQF001 asked whether an *anisotropic resolution field*, a positive definite matrix $A(x)$ attached to every input point and estimated from the local geometry of the data, carries information about uncertainty that local metrics and prototype methods do not already carry, and whether it improves abstention decisions. No manuscript existed; this note reconstructs a precise version of the idea, derives three uncertainty scores from it (a Mahalanobis margin to local prototypes under $A(x)$, the log-volume and the anisotropy of the resolution ellipsoid) and evaluates them as selective classifiers, by risk–coverage curves and their area (AURC), against seven tuned references on 8 datasets with 15 paired folds. The criterion for a “residual gain”, fixed before the run, was a lower AURC than the best reference on a majority of datasets with a paired 95% interval excluding zero. It was not met: the anisotropic field was better than the best reference on 0 datasets and worse, with a 95% bootstrap interval excluding zero, on 6 (5 with a t -interval over folds). Anisotropy improves on the field’s own isotropic ablation on 2–3 datasets, depending on the interval used, and hurts on 2; the evidence that the representation is non-empty is suggestive, not established. An exact statement isolates what anisotropy could add: with the metric equal to the shared within-class inverse covariance, the resolution margin is a strictly increasing function of the LDA posterior margin for two equiprobable classes, equals twice the log-ratio of the two largest posteriors for any number of equiprobable classes, and its departure from LDA under an arbitrary field is $2(x - \bar{m})^\top (A(x) - \Sigma^{-1}) \delta$, so that only the component of the metric along the prototype difference δ can change the score. The formulation is computational throughout; nothing in it refers to physical, let alone quantum, phenomena. With the predefined criterion unmet, the question, in the precise form reconstructed here, is answered in the negative, with numbers.

1 Introduction and scope

A classifier that may abstain trades coverage for accuracy [2, 7, 9]. The trade is governed by a scalar confidence score: examples are accepted in decreasing order of the score and the selective risk is tracked as a function of coverage; the area under this risk–coverage curve (AURC) summarises the quality of the score for a fixed classifier [10, 11]. The simplest scores are the maximum posterior probability of a probabilistic classifier [15], the vote fraction of a k -nearest-neighbour rule [3] and the distance margin of a prototype classifier [18, 23]. Locally adaptive metrics [13, 25] change the geometry in which those distances are measured.

Research line HQF001 proposed to represent uncertainty by an *anisotropic resolution field*: a field of positive definite matrices $A(x)$, one per input point, describing in which directions and at what scale the data around x resolve distinctions between classes, together with an abstention rule derived from it. Its public record stated the objective (“test whether the representation adds information beyond local metrics and prototype-based methods”), the finding (“available comparisons do not establish a residual gain over those references”) and the scope (“the formulation is computational and makes no claim about quantum physical phenomena”). No manuscript existed; the present note reconstructs a precise version of the idea from that record, states one exact result about it, and runs the controlled comparison that the record said was missing, with the success criterion fixed in advance.

Scope. Everything below is a statement about algorithms evaluated on public and synthetic datasets. The word “resolution” names a matrix field; it carries no physical meaning here, and we restate the line’s

own disclaimer: nothing in this note refers to quantum phenomena. Every number in the text is produced by the scripts of Appendix A with a fixed seed and is injected as a generated macro; none is typed. This is version 0.2: version 0.1 (30 September 2026) went through one round of internal review, whose findings are addressed here and listed in the response file of the directory.

2 Setting and definitions

2.1 Selective classification

Let $(x, y) \in \mathbb{R}^d \times \{1, \dots, K\}$ and let h be a classifier with confidence score $g : \mathbb{R}^d \rightarrow \mathbb{R}$ (higher means more confident). For a threshold t , the selective classifier predicts $h(x)$ when $g(x) \geq t$ and abstains otherwise. On a test sample of size n , sorting by decreasing g and accepting the first i points gives coverage $c_i = i/n$ and select risk r_i (the error rate among the accepted points). The risk-coverage curve is $i \mapsto (c_i, r_i)$ and

$$\text{AURC}(g, h) = \frac{1}{n} \sum_{i=1}^n r_i,$$

the discrete version of the area under the curve [11]. Ties in g are frequent for vote-based scores; we compute the *expected* r_i under uniformly random tie-breaking, exactly, block by block (Appendix A). A score with no information about correctness yields, in expectation, r_i equal to the overall error rate at every coverage, so the error rate of h is, in expectation, the AURC of a random ordering and serves as the “no information” level. We also report the selective accuracy $1 - r_i$ at coverages 0.9 and 0.8.

2.2 The anisotropic resolution field

Definition 2.1 (Resolution field and derived scores). Let $D = \{(x_j, y_j)\}_{j=1}^N$ be a training set with standardised features. Fix a neighbourhood size K_m and a shrinkage $\alpha \in (0, 1]$. For a query x , let $N(x)$ be its K_m nearest training points in the Euclidean metric, $C(x)$ their sample covariance, $\tau(x) = \text{tr } C(x)/d$, and

$$S(x) = (1 - \alpha) C(x) + \alpha \tau(x) I, \quad A(x) = S(x)^{-1}.$$

$A(x)$ is the *resolution field*; the ellipsoid $\{u : u^\top A(x) u \leq 1\}$ is the *resolution ellipsoid* at x . For each class c , the *local prototype* $\mu_c(x)$ is the mean of the points of class c in $N(x)$, or the nearest training point of class c when the class is absent from $N(x)$. With $d_c^2(x) = (x - \mu_c(x))^\top A(x) (x - \mu_c(x))$ and $d_{(1)}^2 \leq d_{(2)}^2 \leq \dots$ the sorted values, the field classifier predicts $\arg \min_c d_c^2(x)$ and the three scores are

$$s_{\text{margin}}(x) = d_{(2)}^2(x) - d_{(1)}^2(x), \quad s_{\text{vol}}(x) = -\log \det S(x), \quad s_{\text{anis}}(x) = -\log \frac{\lambda_{\max}(S(x))}{\lambda_{\min}(S(x))}.$$

The margin is the distance to the local decision boundary between the two nearest prototypes, measured in resolution units. The log-volume is a pure density proxy (a large ellipsoid means sparse data, hence low resolution); the anisotropy score ranks elongated neighbourhoods as less certain. Two ablations isolate the geometry: the *isotropic field* sets $\alpha = 1$, so $A(x) = \tau(x)^{-1} I$ and the margin is the Euclidean margin in units of the local scale; the *Euclidean field* sets $A(x) = I$, which is nearest-local-centroid classification. A further variant keeps $A(x)$ but uses the global class means as prototypes. The shrinkage towards a scaled identity is the standard device for neighbourhood covariances in moderate dimension [19]; α and K_m are tuned. Figure 1 shows the estimated ellipsoids on one of the synthetic problems of Section 4.

2.3 References

All references are tuned by inner cross-validation on the same criterion (AURC) and see the same preprocessed features. (a) k -NN with vote-fraction confidence, uniform or distance weights, $k \in \{3, \dots, 31\}$. (b) Prototype methods: nearest class mean (NCM) with Euclidean margin; LDA (prototypes with the global pooled metric $\hat{\Sigma}^{-1}$) and QDA (class-specific metrics), both with maximum-posterior confidence; the LVQ family [17, 18] is represented by NCM. (c) A locally adaptive metric: discriminant adaptive nearest neighbours (DANN) [13], one iteration, $\Sigma(x) = W^{-1/2} [W^{-1/2} B W^{-1/2} + \varepsilon I] W^{-1/2}$ from the within- and between-class scatter of $K_m \in \{50, 100\}$ neighbours, then a k -NN vote under $\Sigma(x)$. (d) Maximum softmax probability of tuned logistic regression and of a random forest [1]. Implementations are from scikit-learn [21] except DANN and the field, which are written in NumPy [12].




Figure 1: Resolution ellipsoids ($K_m = 40$, $\alpha = 0.2$, drawn at 0.6 standard deviations, projected on the two informative coordinates) on a grid of query points of moons-aniso. Along the moons the ellipses follow the local orientation of the data; away from them they inflate.

3 An exact statement: what anisotropy could add

The comparison is only interpretable if one knows what the field score reduces to when the metric is the “right” one. The right one, in the classical setting of Gaussian classes with a shared covariance Σ , is Σ^{-1} : the LDA rule [8, 14]. Write $d_c^2(x) = (x - \mu_c)^\top \Sigma^{-1} (x - \mu_c)$ with the true class means, $s(x) = d_{(2)}^2(x) - d_{(1)}^2(x)$ for the resolution margin under this metric, and

$$p_c(x) = \frac{\pi_c \exp(-d_c^2(x)/2)}{\sum_j \pi_j \exp(-d_j^2(x)/2)}$$

for the LDA posterior with priors π_c , whose two largest values we denote $p_{(1)}(x) \geq p_{(2)}(x)$, attained at classes (1) and (2). The *posterior margin* is $m(x) = p_{(1)}(x) - p_{(2)}(x)$.

Proposition 3.1 (Resolution margin versus LDA posterior margin). *[proved here] Let $K \geq 2$ Gaussian classes share the covariance $\Sigma \succ 0$.*

- (i) *If $K = 2$ and $\pi_1 = \pi_2$, the nearest-prototype rule $\arg \min_c d_c^2$ is the Bayes rule and $m(x) = \tanh(s(x)/4)$. Hence s and m induce the same ordering of every test set, the same risk-coverage curve and the same AUC, and the rejection rules $\{s < t\}$ and $\{m < \tanh(t/4)\}$ coincide.*
- (ii) *If $K = 2$ and $\pi_1 \neq \pi_2$, let $\sigma(x) = d_2^2(x) - d_1^2(x)$ be the signed margin. Then $\log \frac{p_1}{p_2} = \frac{\sigma}{2} + \log \frac{\pi_1}{\pi_2}$ and $m = \tanh(|\sigma + 2 \log(\pi_1/\pi_2)|/4)$. The unsigned score $s = |\sigma|$ is not a function of m , and the nearest-prototype prediction differs from the Bayes prediction on the band $\{-2 \log(\pi_1/\pi_2) < \sigma < 0\}$ (for $\pi_1 > \pi_2$); the prior-corrected score $|\sigma + 2 \log(\pi_1/\pi_2)|$ restores (i).*
- (iii) *If $K \geq 3$ and the priors are equal, $s(x) = 2 \log(p_{(1)}(x)/p_{(2)}(x))$: the resolution margin is a strictly increasing function of the log-ratio of the two largest posteriors. It is not a function of m nor of the maximum posterior $p_{(1)}$. Explicitly, with $\Sigma = I$, $\mu_1 = (0, 0)$, $\mu_2 = (2, 0)$, $\mu_3 = (1, 2)$, the points $x_A = (0.75, 0)$ and $x_B = (0.75, 0.5)$ both have $s = 1$ but $m(x_A) = 0.221$, $m(x_B) = 0.189$ and $p_{(1)}(x_A) = 0.562$, $p_{(1)}(x_B) = 0.481$. For arbitrary priors the identity becomes $s(x) = 2 \log(p_{(1)}/p_{(2)}) - 2 \log(\pi_{(1)}/\pi_{(2)})$ whenever the two nearest prototypes are the two classes of highest posterior, in the same order; without the prior term it fails.*

Proof. (i) With equal priors $\log(p_1/p_2) = (d_2^2 - d_1^2)/2$, so the Bayes rule picks the smaller d_c^2 , and $p_{(1)}/p_{(2)} = e^{s/2}$. Since $p_{(1)} + p_{(2)} = 1$, $m = (e^{s/2} - 1)/(e^{s/2} + 1) = \tanh(s/4)$, a strictly increasing bijection $[0, \infty) \rightarrow [0, 1)$. Every selective-classification quantity depends on g only through the induced ordering (including ties) and on the predictions, which agree. (ii) The first identity is the LDA log-odds; the Bayes prediction is $\text{sign}(\sigma + 2 \log(\pi_1/\pi_2))$, whereas the nearest-prototype prediction is $\text{sign} \sigma$; they differ exactly on the stated band, and two points with $\sigma = \pm a$ have the same $s = a$ but different m unless $\log(\pi_1/\pi_2) = 0$. (iii) In general $\log(p_i/p_j) = \log(\pi_i/\pi_j) + (d_j^2 - d_i^2)/2$. With equal priors the ranking of the posteriors is the reverse ranking of the d_c^2 , so $p_{(1)}/p_{(2)} = e^{s/2}$; with arbitrary priors and the two nearest prototypes equal to classes (1), (2) the same identity gives the prior-corrected form. For the counterexample, $(d_1^2, d_2^2, d_3^2)(x_A) = (0.5625, 1.562, 4.062)$ and $(d_1^2, d_2^2, d_3^2)(x_B) = (0.8125, 1.812, 2.312)$:

in both cases classes 1, 2 are the two nearest with gap 1, but the third distance differs and enters the normalisation of the posteriors. The numbers were also checked by `experiments/lda_identity.py`. $\square$

Proposition 3.2 (Departure from LDA under an arbitrary field). *[proved here] Let $A(x) \succ 0$ be any field, $d_{A,c}^2(x) = (x - \mu_c)^\top A(x)(x - \mu_c)$, and suppose the two nearest prototypes at x under $A(x)$ and under Σ^{-1} are the same pair in the same order, with difference $\delta = \mu_{(1)} - \mu_{(2)}$ and midpoint $\bar{m} = (\mu_{(1)} + \mu_{(2)})/2$. Then*

$$s_A(x) - s(x) = 2(x - \bar{m})^\top (A(x) - \Sigma^{-1}) \delta.$$

In particular: (a) $s_A(x) = s(x)$ whenever $(A(x) - \Sigma^{-1})\delta = 0$, so components of the metric orthogonal to the discriminant direction δ cannot change the score; (b) for the isotropic rescaling $A(x) = \lambda(x)\Sigma^{-1}$ one gets $s_A = \lambda s$, the predictions are unchanged, and the ordering changes as soon as λ is non-constant (two points with $s(x) < s(x')$ swap order when $\lambda(x)/\lambda(x') > s(x')/s(x)$); (c) if the prototypes are replaced by arbitrary points $\nu_c(x)$ (for instance the local prototypes of Definition 2.1), with $\delta_\nu = \nu_{(1)} - \nu_{(2)}$ and $\bar{m}_\nu = (\nu_{(1)} + \nu_{(2)})/2$, and the same pair in the same order is nearest in all three configurations, then

$$s_{A,\nu}(x) - s(x) = 2(x - \bar{m})^\top (A(x) - \Sigma^{-1}) \delta + 2(x - \bar{m})^\top A(x) (\delta_\nu - \delta) + 2(\bar{m} - \bar{m}_\nu)^\top A(x) \delta_\nu.$$

The hypothesis on the order is needed: the identities hold for the signed margins without it, so if the pair is the same but the order is reversed under $A(x)$ they hold with $s_A + s$ on the left and a change of sign on the right; if the pair itself changes, no identity of this form holds.

Proof. For any symmetric M and any two points ν_1, ν_2 , $(x - \nu_2)^\top M(x - \nu_2) - (x - \nu_1)^\top M(x - \nu_1) = 2x^\top M(\nu_1 - \nu_2) - (\nu_1 - \nu_2)^\top M(\nu_1 + \nu_2) = 2(x - \bar{\nu})^\top M(\nu_1 - \nu_2)$. With $M = A(x) - \Sigma^{-1}$ and the true means this is the main identity; (a) is immediate; (b) follows from $M = (\lambda - 1)\Sigma^{-1}$ and $2(x - \bar{m})^\top \Sigma^{-1} \delta = s$; (c) is $s_{A,\nu} - s = [2(x - \bar{m}_\nu)^\top A \delta_\nu - 2(x - \bar{m})^\top A \delta] + [2(x - \bar{m})^\top A \delta - 2(x - \bar{m})^\top \Sigma^{-1} \delta]$ with the first bracket regrouped as stated. The last sentence follows because the signed margins are $2(x - \bar{m})^\top M \delta$ for a fixed labelling of the pair. $\square$

What the experiment can and cannot show. Proposition 3.1(i) says that on shared-covariance data with two balanced classes and the true means as prototypes the field cannot improve on LDA by construction: if $A(x)$ is estimated well it reproduces the LDA ordering, and any difference is estimation error. Proposition 3.2(a) says that, with fixed prototypes, where the field *could* help it must do so through the component of $A(x) - \hat{\Sigma}^{-1}$ along the discriminant direction, which is exactly the quantity that class-dependent covariances (QDA regime) or curved boundaries make non-zero. The method as run uses *local* prototypes, and part (c) shows that two further terms then enter, which (a) does not control; the statement about “the component along δ ” is therefore exact for fixed prototypes and a heuristic for the method as run. The variant “aniso, global prototypes” of Section 5 isolates the prototype terms empirically, and the control synth-lda measures estimation error and prototype effects together without separating them. This dictates the synthetic datasets: one shared-covariance control where no gain is possible, one class-dependent-covariance problem and one curved-boundary problem with anisotropic noise where a gain is possible in principle, and one generic problem with redundant features. *[proved here for fixed prototypes / heuristic for local prototypes / design]* A second heuristic explains the direction of the results: $C(x)$ is the covariance of the *mixture* of classes in $N(x)$; where two classes meet it equals the within-class covariance plus a between-class term proportional to $\delta\delta^\top$, so $A(x)$ is *smaller* along δ precisely near boundaries, shrinking the margin there. The direction of this bias is consistent with abstention, but its size is set by the local class balance rather than by the class geometry; DANN separates the within- and between-class parts and avoids it. *[heuristic]*

The numerical verification (`experiments/lda_identity.py`, true parameters, $d = 5$, $n = 4000$ points per case) gives: for $K = 2$ and equal priors, Spearman correlation between s and m equal to 1.000000, $\max |m - \tanh(s/4)| = 2.8 \times 10^{-16}$, AURC difference 0; for priors (0.8, 0.2), Spearman 0.9725, monotone: no, predictions differ on 0.57% of the points, and the corrected score is monotone (yes, deviation 4.4×10^{-16}); for $K = 3$ with equal priors (error rate 11.2%), $\max |s - 2 \log(p_{(1)}/p_{(2)})| = 7.1 \times 10^{-15}$, monotone in the log-ratio: yes, in m : no, in $p_{(1)}$: no, with AURC $\times 100$ of 2.415 (s), 2.392 (m) and 2.380 ($p_{(1)}$): the three orderings are different but nearly equivalent in this instance (Spearman 0.9993 and 0.9971); for $K = 3$ with priors (0.6, 0.3, 0.1) the uncorrected identity fails ($\max |s - 2 \log(p_{(1)}/p_{(2)})| = 3.58$) while the prior-corrected form holds to 7.1×10^{-15} on the 85.3% of points where the nearest pair is the top-posterior pair; the decomposition of Proposition 3.2(c) has relative residual 4.0×10^{-14} over 500 random fields, means and prototypes. Figure 2 shows the first and third cases.




Figure 2: Proposition 3.1 on samples from the true model. Left: $K = 2$, equal priors, the posterior margin is exactly $\tanh(s/4)$. Right: $K = 3$, equal priors, neither the posterior difference nor the maximum posterior is a function of s .

dataset	n	d	K	Bayes err. (%)	description
iris	150	4	3	–	Fisher’s iris
wine	178	13	3	–	UCI wine
breast-cancer	569	30	2	–	Wisconsin diagnostic
digits	1797	20	10	–	UCI optical digits, PCA to 20 comp.
synth-informative	1200	10	2	–	make_classification : 5 inf. + 3 red. + 2 noise, 2 clusters/class, 3% label noise
moons-aniso	1200	5	2	–	two moons + anisotropic noise (sd 0.30×0.06 , 30°) + 3 nuisance dims
synth-classcov	1200	6	2	9.9	two Gaussians, $d = 6$, rotated class-specific covariances (spectrum $2 \rightarrow 0.25$), mean shift 1.54
synth-lda	1200	6	3	9.8	three Gaussians, $d = 6$, one shared covariance (spectrum $2 \rightarrow 0.05$), mean shift 1.20

Table 1: Datasets. d is the dimension after preprocessing; the Bayes error is known only for the Gaussian problems.

4 Experimental protocol and predefined criterion

Datasets. Table 1. The real datasets are those shipped with scikit-learn [8, 21, 22]; digits is reduced to 20 principal components (fitted on each training fold) so that neighbourhood covariances are estimable for all local methods alike. The synthetic problems follow the design paragraph of Section 3: *synth-informative* is **make_classification** with informative, redundant and noise features; *moons-aniso* adds anisotropic Gaussian noise (standard deviations 0.30 and 0.06 along axes rotated by 30°) and three nuisance dimensions to the two-moons problem; *synth-classcov* has two Gaussian classes in $d = 6$ with the same covariance spectrum ($2 \rightarrow 0.25$, geometric) under two different random rotations; *synth-lda* has three Gaussian classes with one shared covariance (spectrum $2 \rightarrow 0.05$). In the last two the mean separation is chosen by bisection on a Monte-Carlo estimate of the Bayes error, with target 10% and search interval $[0.05, 20]$; the bisection now reports whether it ended at a bound (it did not: shifts 1.54 and 1.20, Bayes errors 9.9% and 9.8%). In version 0.1 the *synth-classcov* spectrum was $(2, 1, 0.5, 0.25, 0.1, 0.05)$, for which the rotated covariances alone give a 7.7% Bayes error; the bisection stopped at its lower bound, the class means coincided, and the linear references were at chance level there, so the field’s wins over them were a consequence of the design. The dataset was regenerated for version 0.2 (same rotations and draws) after the v0.1 results had been seen; the v0.1 numbers are reported in Section 5 and kept in `results/results_v01.json`, and the criterion verdict is the same under either version.

Evaluation. Repeated stratified 5-fold cross-validation with 3 repeats (15 outer folds, identical for all methods). Inside each training fold, features are standardised and every method’s hyper-parameters are chosen by 3-fold inner cross-validation minimising AUC; the chosen configuration is refitted on the whole training fold and scored on the test fold. Paired per-fold differences of AUC are summarised by their mean and three 95% intervals: a percentile bootstrap over folds ($B = 20\,000$) [6], which is the interval named

dataset	references							field variants						
	k -NN	NCM	LDA	QDA	LogReg	RForest	DANN	err.	aniso	iso	Eucl.	vol.	anis.	glob. prot.
iris	0.68	2.73	0.17	0.23	0.25	0.54	0.19	2.0	0.37	0.65	0.76	1.54	6.79	1.77
wine	0.69	0.57	0.10	0.03	0.11	0.15	0.70	0.9	0.08	0.28	0.34	0.65	0.64	1.10
breast-cancer	0.57	1.15	0.61	0.54	0.25	0.64	0.86	2.3	0.83	0.45	0.57	4.73	4.35	1.66
digits	0.28	2.97	1.54	1.80	0.64	0.41	0.25	2.5	0.26	1.04	4.11	0.66	1.42	2.16
synth-informative	4.81	21.04	18.00	6.41	17.54	3.51	3.96	10.5	4.81	4.95	4.95	17.10	15.60	21.31
moons-aniso	0.82	2.36	2.18	2.23	2.17	0.62	1.23	4.0	1.19	1.23	1.64	4.67	10.18	2.00
synth-classcov	2.55	8.51	6.30	1.69	5.96	3.42	2.83	9.6	2.77	2.30	2.31	16.70	15.10	10.23
synth-lda	2.90	5.05	1.58	1.64	1.63	3.67	2.26	8.9	2.83	2.94	3.10	15.02	14.42	4.48

Table 2: AURC $\times 100$, mean over 15 folds (lower is better; best reference in bold). “err.” is the error rate $\times 100$ of the best reference at full coverage, i.e. the expected AURC of an uninformative score for that classifier. “aniso” is the full method of Definition 2.1; “iso” and “Eucl.” are its ablations; “vol.” and “anis.” rank by geometry alone (with the predictions of the full method); “glob. prot.” keeps $A(x)$ with the global class means as prototypes.

in the predefined criterion; a Student t interval over folds ($15 - 1$ degrees of freedom); and, as an extreme, the t interval with the Nadeau–Bengio corrected variance $(1/J + \rho)s^2$, $\rho = n_{\text{test}}/n_{\text{train}} = 0.25$ [20]. We also report wins/ties/losses over folds and a two-sided Wilcoxon signed-rank p -value [26] as a descriptive secondary check. Folds of a repeated cross-validation are not independent (each test point appears 3 times; training sets overlap by about 80%), so the bootstrap and t intervals are optimistic [5, 20]. This optimism cannot create a “better” verdict where there is none, so the criterion count is conservative; it does inflate every “worse” verdict and every positive ablation count below, which we therefore also report with the t interval and with the Nadeau–Bengio correction. No multiplicity correction is applied to the 56 intervals of Table 4 and the 48 of Table 5; those counts are descriptive, and we do not perform the cross-dataset tests of Demšar [4]. An interval whose end lies within 0.02 (AURC $\times 100$) of zero is marked as *borderline* ($^\circ$) and should be read as inconclusive. On iris and wine the test folds hold 30 and 35 points and the best reference has AURC exactly 0 in 8 and 10 of 15 folds, so exact ties are frequent and the verdicts there rest on differences of one or two errors.

Predefined criterion. Fixed before the run (30 September 2026; the dated record is CRITERIO_HQF001.md in the directory, written for version 0.2, which states explicitly that no independent timestamp earlier than the run exists beyond the continuity note): the field score shows a *residual gain* if its AURC is lower than that of the best reference (the reference with the lowest mean AURC on that dataset, among the 7 references) on a majority (≥ 5 of 8) of datasets, with the paired 95% bootstrap interval excluding zero. Choosing the best reference per dataset on the same folds favours the references slightly; we therefore also report the field against each reference separately. The hyper-parameter key of the field was corrected between the smoke test and the reference run (it now carries the nominal rather than the capped K_m , so that inner and outer configurations match on small folds); no reference result had been seen. Version 0.2 re-ran the complete benchmark with the current script, so the stored metadata and the code coincide. Total compute: 285 s of CPU (4.7 min) for the benchmark, single-threaded, plus 0.04 s for the verification of Section 3.

5 Results

5.1 Main comparison

Tables 2 and 3 give the main result. Against the best reference, the anisotropic field is better with a 95% bootstrap interval excluding zero on 0 of 8 datasets (none), worse on 6 (iris, breast-cancer, synth-informative, moons-aniso, synth-classcov, synth-lda) and inconclusive on 2 (wine, digits); with the t interval the counts are 0/5/3 (inconclusive: iris, wine, digits), and with the Nadeau–Bengio correction 0/3/5. The predefined criterion is therefore **not met** under every reading. The closest the field comes to the best reference is on digits ($\Delta = +0.01$); the largest loss is on synth-informative ($\Delta = +1.30$). The isotropic and Euclidean ablations are better than the best reference on 0 and 0 datasets and worse on 8 and 8 (8 and 8 with the t interval); the geometry-only scores are worse on 8 and 8 (full tables in results/tables.md).

Table 4 removes the best-reference selection. Counting datasets on which the field is better/worse with the bootstrap interval, and in parentheses with the t interval: k -NN 1/3 (0/1), NCM 7/0 (7/0),

dataset	best ref.	Δ	bootstrap 95%	t 95%	NB 95%	W/T/L	p_W
iris	LDA	+0.20	[+0.05, +0.43]	[-0.03, +0.43]	[-0.30, +0.70]	1/6/8	0.028
wine	QDA	+0.05 ^o	[-0.005, +0.10]	[-0.01, +0.11]	[-0.08, +0.18]	4/4/7	0.129
breast-cancer	LogReg	+0.58	[+0.34, +0.89]	[+0.26, +0.90]	[-0.11, +1.27]	0/0/15	0.000
digits	DANN	+0.01	[-0.11, +0.14]	[-0.13, +0.15]	[-0.30, +0.32]	7/0/8	0.934
synth-informative	RForest	+1.30	[+0.78, +1.86]	[+0.69, +1.92]	[-0.03, +2.64]	2/0/13	0.001
moons-aniso	RForest	+0.56	[+0.35, +0.75]	[+0.33, +0.80]	[+0.06, +1.07]	2/0/13	0.000
synth-classcov	QDA	+1.08	[+0.86, +1.33]	[+0.81, +1.35]	[+0.49, +1.67]	0/0/15	0.000
synth-lda	LDA	+1.25	[+1.04, +1.50]	[+0.98, +1.52]	[+0.67, +1.83]	0/0/15	0.000

Table 3: Anisotropic field against the best reference: paired difference $\Delta = \text{AURC}(\text{field}) - \text{AURC}(\text{best reference})$, $\times 100$ (negative favours the field), with three 95% intervals over folds (percentile bootstrap, the predefined one; Student t ; NB: t with the Nadeau–Bengio correction), wins/ties/losses of the field over the 15 folds and the Wilcoxon p -value. ^o: borderline (an end of the bootstrap or t interval within 0.02 of zero). The same table for the isotropic and Euclidean ablations is in `results/tables.md`.

dataset	vs k -NN	vs NCM	vs LDA	vs QDA	vs LogReg	vs RForest	vs DANN
iris	-0.31	-2.36	+0.20 ^b	+0.14 ^{bo}	+0.11	-0.17	+0.17 ^b
wine	-0.61 ^b	-0.49	-0.02	+0.05 ^o	-0.03	-0.07	-0.61 ^b
breast-cancer	+0.25 ^{bo}	-0.32 ^o	+0.21	+0.29	+0.58	+0.18	-0.04
digits	-0.02	-2.71	-1.28	-1.54	-0.38	-0.15	+0.01
synth-informative	+0.01	-16.23	-13.18	-1.60	-12.72	+1.30	+0.86
moons-aniso	+0.37	-1.18	-0.99	-1.05	-0.99	+0.56	-0.05
synth-classcov	+0.22 ^{bo}	-5.73	-3.53	+1.08	-3.18	-0.65	-0.06
synth-lda	-0.07	-2.22	+1.25	+1.19	+1.20	-0.84	+0.56

Table 4: Anisotropic field against each reference separately: mean paired Δ of $\text{AURC} \times 100$ (negative favours the field). Bold: field better, with both the bootstrap and the t interval excluding zero; italic: field worse, both intervals; ^b (^t): only the bootstrap (only the t) interval excludes zero; ^o: borderline; roman without mark: inconclusive. Full intervals in `results/tables.md`.

LDA 4/3 (4/2), QDA 3/3 (3/2), logistic regression 4/2 (4/2), random forest 3/2 (3/2), DANN 1/3 (0/2). 5 of these 56 intervals are borderline (iris vs QDA; wine vs QDA; breast-cancer vs k -NN; breast-cancer vs NCM; synth-classcov vs k -NN) and 7 are decided by the bootstrap interval only (iris vs LDA; iris vs QDA; iris vs DANN; wine vs k -NN; wine vs DANN; breast-cancer vs k -NN; synth-classcov vs k -NN); all of them should be read as inconclusive. The regime-specific picture is as follows. The field beats the global-prototype methods on the datasets where their linear boundaries are wrong (with both intervals: LDA on digits, synth-informative, moons-aniso, synth-classcov; logistic regression on digits, synth-informative, moons-aniso, synth-classcov; NCM on all but breast-cancer), and on synth-classcov, the class-dependent-covariance problem, it beats the random forest (-0.65 [-0.94, -0.32]), is inconclusive against DANN (-0.06 [-0.29, +0.12]) and is worse than k -NN with the bootstrap interval (+0.22 [+0.03, +0.43], borderline; t interval [-0.003, +0.45]). On every dataset at least one reference is at least as good, and on synth-classcov that reference is QDA ($\text{AURC} \times 100$ 1.69 against 2.77 for the field), whose class-specific metric is the population version of what the field estimates locally. Moreover, on synth-classcov anisotropy is a handicap rather than a help: the isotropic and Euclidean ablations reach 2.30 and 2.31 (Table 2; aniso–iso +0.47 [+0.35, +0.58], Table 5), so whatever the field has to offer in this regime is the local-prototype margin, not the metric. For the record, in version 0.1 (coinciding means, Bayes error 7.7%) QDA was also best (0.97 against 1.40; $\Delta = +0.43$, 0/0/15 wins/ties/losses), the Euclidean ablation (1.23) again beat the full method, the wins over NCM, LDA and logistic regression ($\Delta \approx -47.04$) were a consequence of the design (those references were at chance level, error 44.8%), and of the wins over k -NN, DANN and the random forest (-0.29, -0.37, -1.08) only the last persists in version 0.2.

5.2 Ablations: does anisotropy add information?

Table 5 answers the line’s question in its narrow form. Allowing anisotropy improves on the isotropic field (same neighbourhoods, same prototypes, scalar metric) with a bootstrap interval excluding zero on 3 datasets (iris, wine, digits) and with a t interval on 2 (wine, digits); with the Nadeau–Bengio correction on 0. It hurts on 2 (breast-cancer, synth-classcov) with the bootstrap and on 2 with the t interval. The fold counts behind the positive cases are digits 15/0/0, wine 8/5/2, iris 8/4/3 (wins/ties/losses). Against the

dataset	aniso-iso				W/T/L	aniso-Eucl.	Eucl.-NCM	aniso-vol.	aniso-anis.	aniso-glob.
	Δ	bootstrap 95%	t 95%							
iris	-0.28°	[-0.56, -0.05]	[-0.58, +0.01]	8/4/3	-0.40	-1.97	-1.18	-6.42	-1.40	
wine	-0.20°	[-0.37, -0.05]	[-0.39, -0.01]	8/5/2	-0.26 ^{b°}	-0.23	-0.56	-0.56 ^b	-1.02	
breast-cancer	+0.37	[+0.19, +0.60]	[+0.14, +0.61]	1/0/14	+0.25°	-0.57	-3.90	-3.53	-0.84	
digits	-0.78	[-1.22, -0.41]	[-1.25, -0.32]	15/0/0	-3.85	+1.14	-0.40	-1.16	-1.91	
synth-informative	-0.14	[-0.51, +0.19]	[-0.54, +0.26]	8/0/7	-0.14	-16.09	-12.28	-10.79	-16.49	
moons-aniso	-0.05°	[-0.10, +0.005]	[-0.11, +0.01]	10/0/5	-0.46	-0.72	-3.49	-8.99	-0.81	
synth-classcov	+0.47	[+0.35, +0.58]	[+0.34, +0.60]	0/0/15	+0.47	-6.20	-13.93	-12.32	-7.46	
synth-lda	-0.11	[-0.37, +0.11]	[-0.38, +0.16]	6/0/9	-0.27	-1.95	-12.19	-11.59	-1.66	

Table 5: Ablations: mean paired difference of $\text{AURC} \times 100$ (first minus second; negative favours the first). For aniso-iso, the two 95% intervals and the wins/ties/losses of the anisotropic field over folds are shown; for the other pairs, bold (italic) marks a difference whose bootstrap and t intervals both exclude zero in favour of the first (second) member, ^b and ^t a difference decided by one interval only, ° a borderline interval. Full intervals in `results/tables.md`; direct comparisons with the references are in Table 4.

Euclidean field the counts are 4 and 1 (bootstrap; 3 and 1 with t). So anisotropy improves on the field’s own isotropic ablation on 2–3 datasets depending on the interval used and hurts on 2: the evidence that the anisotropic metric carries information about correctness that a local scale alone does not is suggestive, not established. Replacing the local prototypes by the global class means while keeping $A(x)$ is worse on 8 datasets with both intervals: the local metric is meaningful only at the scale of the neighbourhood from which it was estimated. 5 of the 48 ablation intervals are borderline (iris (aniso-iso); wine (aniso-iso); moons-aniso (aniso-iso); wine (aniso-Eucl.); breast-cancer (aniso-Eucl.)).

5.3 Geometry-only scores, fixed coverage, selected hyper-parameters

The geometry-only scores (section “Geometry-only scores” of `results/tables.md`) are far from the margin. Ranking by log-volume is *worse* than a random ordering of the same predictions on 3 of 8 datasets (synth-informative, synth-classcov, synth-lda) and above half of the random level on 5; ranking by anisotropy is worse than random on 5 (iris, synth-informative, moons-aniso, synth-classcov, synth-lda). The margin is below half of the random level on every dataset. The information the field carries about correctness is therefore in the margin, i.e. in *where the prototypes are* relative to the ellipsoid, and not in the ellipsoid’s size or shape by itself.

The accuracies at fixed coverage (`results/tables.md`; mean risk-coverage curves in `figures/fig_rc_curves.pdf`) show the same picture. On 6 datasets the field’s classifier is already less accurate than the best reference at full coverage, by 1.8 (iris), 4.4 (breast-cancer), 3.0 (synth-informative), 4.9 (moons-aniso), 1.8 (synth-classcov), 2.9 (synth-lda) percentage points, so part of the AURC gap is a weaker classifier rather than a weaker ranking. On digits the field is slightly *more* accurate than DANN at full coverage (97.7% against 97.5%) yet behind at 90% coverage (99.2% against 99.7%), so there the ranking is the weaker part; the two effects cancel in the AURC. The hyper-parameters selected for the anisotropic field over all 120 dataset-fold pairs were $\alpha = 0.05$ in 9% of the cases, $\alpha = 0.2$ in 18% and $\alpha = 0.5$ in 72%; $K_m = 20$ in 31%, 40 in 16% and 80 in 53%: the tuning prefers the strongest shrinkage on offer, i.e. a mildly anisotropic field.

6 What can be claimed

Table 6 lists the claims with their status. The main one is negative and is stated as such: under a criterion fixed in advance, on 8 datasets chosen to include the regimes where a local anisotropic metric could help, the resolution-field score does not deliver a residual gain over tuned local-metric, prototype and softmax references, under any of the three interval readings. The comparison with the best reference is inconclusive only on wine, digits with the bootstrap interval (wine: 4/4/7 wins/ties/losses, digits: 7/0/8) and additionally on iris with the t interval (1/6/8); nowhere does the field win. The positive finding is narrow and interval-dependent: the anisotropic metric ranks correctness better than an isotropic local scale on 2–3 datasets and worse on 2, so the representation may be non-empty, but it is redundant given the references.

claim	status
LDA posterior and Bayes optimality for shared-covariance Gaussians	classical
$K = 2$, equal priors: the LDA posterior margin equals $\tanh(s/4)$ of the resolution margin under Σ^{-1} ; identical risk-coverage curves (Prop. 3.1(i))	proved here
Unequal priors: identity fails for the unsigned margin; restored by the prior-corrected signed margin (Prop. 3.1(ii))	proved here
$K \geq 3$, equal priors: margin equals twice the top-two posterior log-ratio, not a function of the posterior difference or of the maximum posterior; with unequal priors a prior term is needed (Prop. 3.1(iii))	proved here, examples verified
Departure from LDA under any field is $2(x - \bar{m})^\top (A - \Sigma^{-1})\delta$ with fixed prototypes, plus two terms with arbitrary ones (Prop. 3.2)	proved here
“Where the field could help it must do so along δ ” for the method as run (local prototypes)	heuristic
Unsupervised local covariance is inflated along δ near boundaries	heuristic
Predefined criterion for a residual gain of the field over the best reference: not met (0 better, 6 worse with the bootstrap interval, 5 with the t interval, of 8)	computationally verified
Anisotropy improves on the isotropic ablation on 2–3 datasets depending on the interval and hurts on 2; the representation being non-empty is suggestive, not established	computationally verified, interval-dependent
On synth-classcov (class-dependent covariances, version 0.2) the field beats the linear references and the random forest but not QDA (worse), k -NN or DANN (worse or inconclusive), and its isotropic and Euclidean ablations beat the full method: what helps there is the local-prototype margin, not anisotropy	computationally verified, regime-specific; dataset regenerated after v0.1
Geometry-only scores (volume, anisotropy) are close to uninformative	computationally verified
Behaviour on other data, in high dimension, or with supervised field estimation	not addressed
Any statement about physical or quantum phenomena	not made

Table 6: Claims and their status. “Computationally verified” means reproduced by the scripts of Appendix A with the fixed seed on the listed datasets; it is not a proof and does not generalise beyond them.

Limitations. The field is one specific reconstruction of an idea for which no manuscript existed; other formulations (a supervised field from the Fisher information of a local probabilistic model, a field learned end-to-end, or a combination of the margin with the log-volume) were not run, and the second and third could in principle change the outcome. The datasets are small to moderate ($n \leq 1797$, $d \leq 30$) and low-dimensional after preprocessing; neighbourhood covariances in higher dimension would need stronger regularisation and the comparison could shift either way. Inner tuning on AURC with 3 folds is noisy on the smallest datasets. The bootstrap and t intervals over repeated folds are optimistic and the best reference is selected on the same folds; the first favours every “worse” and every positive ablation verdict (which is why the Nadeau–Bengio reading is also given), the second favours the references; neither can produce a “better” verdict for the field. DANN was run for one iteration with $\varepsilon = 1$ as in the original paper and not tuned further. QDA at `reg_param = 0` was excluded because it fails on rank-deficient inner folds. The synth-classcov problem was regenerated between versions 0.1 and 0.2 (Section 4).

Next steps for the research line. If the line is reopened at all, the directions consistent with the results above are: (a) a supervised field, $A(x)$ built from within-class rather than mixture covariances (the DANN decomposition) with the field margin rather than a vote as the score, tested on synth-classcov and moons-aniso first, since these are the regimes where Proposition 3.2 allows a gain; (b) the margin combined with the log-volume as a two-dimensional rejection rule, evaluated under the same criterion; (c) a higher-dimensional benchmark with explicit covariance regularisation. Absent these, the line is proposed to close for the unsupervised formulation evaluated here: on the problems studied, the field is a valid but redundant representation of uncertainty for abstention.

A Reproducibility

The directory contains `experiments/selective_benchmark.py` (benchmark; `--fast` for a reduced run, `--resummarise` to recompute the statistics from stored per-fold results), `experiments/lda_identity.py` (verification of Section 3), `experiments/make_figures.py` and `experiments/make_numbers.py`, which converts `results/results.json` and `results/results_identity.json` into the macros and table bodies of this document; `results/results_v01.json` is the version 0.1 run. Script v0.2 (2026-10-03); seed 20260930; Python 3.11.15, NumPy 2.4.6, SciPy 1.17.1 [24], scikit-learn 1.9.1, Matplotlib [16]; BLAS threads limited to one; 285 s of CPU for the benchmark (wall 287 s).

Tie handling. Sort by decreasing score and group tied scores into blocks; within a block of size b with e errors, the expected cumulative error after accepting $j \leq b$ of its points is ej/b , and r_i is the expected cumulative error divided by i . This is the exact expectation of the risk–coverage curve under uniform random tie-breaking; the AURC is its mean over i and the accuracy at coverage c is $1 - r_{\lceil cn \rceil}$.

Preprocessing and inner selection. Standardisation (and, for digits, the 20-component PCA) is fitted on each outer training fold and applied to its test fold, so no test information reaches any method. The inner 3-fold cross-validation then splits that already standardised training fold: the inner validation folds have seen their own statistics through the scaler, which can only affect hyper-parameter selection, not the test numbers. Inner selection minimises the mean inner AURC; the hyper-parameter key carries the nominal K_m so that inner and outer configurations match on small folds; a configuration that fails on some inner fold (QDA with a singular class covariance) is discarded rather than averaged over fewer folds: 0 configurations were discarded in the reference run (none). On iris, $K_m = 80$ with 120 training points means that the inner “local” covariance (80 of 96 points) is nearly the global one; this change of regime between inner and outer folds is inherent to the sample size. The hyper-parameter grids are listed in `results/tables.md` and in the `grids` field of `results/results.json`; the synthetic generators are in `make_datasets`, and the class-mean separation of synth-classcov and synth-lda is found by 40 bisection steps on a Monte-Carlo estimate (40 000 points) of the Bayes error.

Supplementary tables. `results/tables.md` (generated with the JSON) holds, with the same conventions as Table 3: every field variant against the best reference and the anisotropic field against each reference (three intervals, wins/ties/losses, Wilcoxon p , borderline flags); all ablation pairs; the geometry-only scores against the error rate of their own tuned classifier; the accuracies at coverage 100%, 90% and 80%; the hyper-parameter grids and the configurations selected per fold. The mean risk–coverage

curves of the anisotropic field, its isotropic ablation and the best reference on digits, moons-aniso and synth-classcov are in `figures/fig_rc_curves.pdf`.

References

- [1] Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, 2001.
- [2] C. K. Chow. On optimum recognition error and reject tradeoff. *IEEE Transactions on Information Theory*, 16(1):41–46, 1970.
- [3] Thomas Cover and Peter Hart. Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1):21–27, 1967.
- [4] Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006.
- [5] Thomas G. Dietterich. Approximate statistical tests for comparing supervised classification learning algorithms. *Neural Computation*, 10(7):1895–1923, 1998.
- [6] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman & Hall, New York, 1993.
- [7] Ran El-Yaniv and Yair Wiener. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11:1605–1641, 2010.
- [8] Ronald A. Fisher. The use of multiple measurements in taxonomic problems. *Annals of Eugenics*, 7(2):179–188, 1936.
- [9] Vojtěch Franc, Daniel Průša, and Václav Voráček. Optimal strategies for reject option classifiers. *Journal of Machine Learning Research*, 24(11):1–49, 2023.
- [10] Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems 30*, 2017.
- [11] Yonatan Geifman, Guy Uziel, and Ran El-Yaniv. Bias-reduced uncertainty estimation for deep neural classifiers. In *International Conference on Learning Representations*, 2019.
- [12] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, et al. Array programming with NumPy. *Nature*, 585:357–362, 2020.
- [13] Trevor Hastie and Robert Tibshirani. Discriminant adaptive nearest neighbor classification. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 18(6):607–616, 1996.
- [14] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer, New York, 2nd edition, 2009.
- [15] Dan Hendrycks and Kevin Gimpel. A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations*, 2017.
- [16] John D. Hunter. Matplotlib: A 2D graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007.
- [17] Teuvo Kohonen. The self-organizing map. *Proceedings of the IEEE*, 78(9):1464–1480, 1990.
- [18] Teuvo Kohonen. *Self-Organizing Maps*, volume 30 of *Springer Series in Information Sciences*. Springer, Berlin, 1995.
- [19] Olivier Ledoit and Michael Wolf. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2):365–411, 2004.
- [20] Claude Nadeau and Yoshua Bengio. Inference for the generalization error. *Machine Learning*, 52(3):239–281, 2003.
- [21] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [22] W. Nick Street, William H. Wolberg, and Olvi L. Mangasarian. Nuclear feature extraction for breast tumor diagnosis. In *Biomedical Image Processing and Biomedical Visualization (IS&T/SPIE 1993)*, volume 1905, pages 861–870, 1993.
- [23] Robert Tibshirani, Trevor Hastie, Balasubramanian Narasimhan, and Gilbert Chu. Diagnosis of multiple cancer types by shrunken centroids of gene expression. *Proceedings of the National Academy of Sciences*, 99(10):6567–6572, 2002.
- [24] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, et al. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nature Methods*, 17:261–272, 2020.
- [25] Kilian Q. Weinberger and Lawrence K. Saul. Distance metric learning for large margin nearest neighbor classification. *Journal of Machine Learning Research*, 10:207–244, 2009.
- [26] Frank Wilcoxon. Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83, 1945.