

# Corrections with Safe Reversion Across Estimation Tasks:

What a Held-Out Check Can and Cannot Guarantee

Luciano Silva Alarco

Pontificia Universidad Católica del Perú

ORCID 0000-0002-3395-3129

Working draft v0.3 (two rounds of internal review applied) — research line OMR001 — 3 October 2026

## Abstract

Research line OMR001 proposed post-estimation *correction operators* with *reversion to a safe reference* when a held-out check fails, and recorded that no sufficient differentiation from existing constructs had been shown. This note reconstructs it in the Gaussian location model and proves what it can and cannot guarantee. (i) No correction fitted on sources whose law is free of the target parameter improves the minimax risk; for  $d \leq 2$  every rule that ever corrects is worse than the reference somewhere. (ii) With a one-sided level- $\alpha$  check on  $m$  held-out observations, the excess risk over the reference is at most  $\min\{\alpha\mathbb{E}[\Delta^+], \kappa(\alpha)(2\sigma/\sqrt{m})\mathbb{E}\|C - R\|\}$ , with  $\kappa(\alpha)$  the best constant, and at most a cap uniform in the parameter, attained in order  $m^{-1/2}$  for a fixed estimation sample; the probability that the correction is used and is harmful is at most  $\alpha\mathbb{P}(\Delta > 0)$ . (iii) The reverting estimator *is* penalised hold-out selection between two candidates: its guarantee is a selection guarantee, indifferent to the operator. A fixed-seed simulation confirms every bound; under a predefined criterion the guaranteed construction is useful only when tasks are close (between-task variance at most 2 times the estimation variance with plain hold-out), because the split costs 25.0% of the full-data reference risk; the variants that avoid the split are classical constructs without a guarantee proved here. Nothing is claimed about general accuracy or transfer guarantees.

## 1 Introduction

Suppose an analyst estimates the same kind of quantity, a mean or a regression coefficient, in a family of related tasks. Each task has its own *reference* estimator, computed from that task’s data alone and trusted because it is unbiased or minimax. A *correction operator* maps the reference estimate to a corrected one using what was learned on the other tasks: shrinkage toward a pooled centre, an additive bias adjustment, a linear map fitted on source tasks. Since the correction can hurt when the target task does not share the structure, the proposal of research line OMR001 was to *revert* to the reference whenever a held-out check on the target task does not favour the correction. The public record of the line states its objective (evaluate corrections on held-out tasks and establish conditions for transfer), its finding (sufficient differentiation from existing model combinations and operator-based corrections has not been demonstrated) and its scope (no general accuracy improvement or transfer guarantee is claimed).

The original definitions of the operators and of the “safe reference” are not available. This note therefore reconstructs a self-contained version of the construction (Section 2) and asks the question the line left open: *what, exactly, can a correction with safe reversion guarantee, and what can it not?* The answers are elementary but, we think, worth having in one place with proofs and with their constants.

Section 3 proves what *cannot* be guaranteed (any gain is a statement about average risk over tasks), and Section 4 what *can*: a do-no-harm bound of order  $\sigma\mathbb{E}\|C - R\|/\sqrt{m}$  with the best constant for one-sided checks at a fixed level, a cap uniform in the parameter, of exact order  $m^{-1/2}$  for a fixed estimation sample, a bound on harmful transfer events, and near-oracle selection. Section 5 shows that the reverting estimator is penalised hold-out selection between two candidates, so that its safety is a property of selection and not of the operator, which explains why no differentiation should have been expected from the safety mechanism itself. Section 6 checks the bounds and reports, under a predefined criterion, where “useful transfer” occurs.

Every mathematical statement below carries a status tag in Table 4: proved here, classical, from the literature, computationally verified, or empirical only. Every number in the text is generated from the results of one script with a fixed seed (Appendix A).

## 2 Setting and definitions

**Tasks and data.** A task is a parameter  $\theta \in \mathbb{R}^d$  together with observations whose law depends on  $\theta$ . We work throughout in the Gaussian location model with known variance: an observation of task  $\theta$  is  $Y \sim N(\theta, \sigma^2 I_d)$ . The *target* task has parameter  $\theta$  and provides  $n = n_e + m$  independent observations; the first  $n_e$  form the *estimation sample*, with mean  $\bar{X}_e \sim N(\theta, \sigma_e^2 I_d)$ ,  $\sigma_e^2 = \sigma^2/n_e$ , and the last  $m$  form the *held-out sample*  $Y_1, \dots, Y_m$ , with mean  $\bar{Y} \sim N(\theta, \sigma^2 I_d/m)$ . *Source* tasks  $s = 1, \dots, S$  have parameters  $\theta_s$  and produce data  $D$  whose law may or may not carry information about  $\theta$ ; this is the whole question. We write  $\mathcal{F}$  for the  $\sigma$ -field generated by the estimation sample and by  $D$ ; the held-out sample is independent of  $\mathcal{F}$  given  $\theta$ .

**Loss and risk.** For an estimate  $\delta$  of  $\theta$  the loss is  $\ell(\delta) = \|\delta - \theta\|^2$  and the risk is  $\mathcal{R}(\theta, \delta) = \mathbb{E}_\theta \ell(\delta)$ . When  $\theta$  is itself drawn from a distribution over tasks, the average of the risk is the Bayes risk.

**Definition 2.1** (Reference, correction, reversion). (a) The *reference* estimator is the task-specific mean  $R = \bar{X}_e$ ; it is unbiased and minimax for  $\theta \in \mathbb{R}^d$  (Theorem 3.1(a)).

(b) A *correction operator* is an  $\mathcal{F}$ -measurable map producing a corrected estimate  $C = C(R; D) \in \mathbb{R}^d$ . Three instances are used below: the affine shrinkage  $C = \hat{\mu} + \lambda(R - \hat{\mu})$  toward a centre  $\hat{\mu}$  estimated from the sources, with  $\lambda$  either estimated from the sources (the empirical Bayes rule of Efron and Morris [6, 7], Morris [15]), fixed in advance, or equal to zero (full pooling).

(c) The *held-out check* compares the held-out losses of the two candidates,

$$\hat{D} = \|C - \bar{Y}\|^2 - \|R - \bar{Y}\|^2,$$

and the *reverting estimator* with margin  $c \geq 0$  ( $\mathcal{F}$ -measurable) is

$$\hat{\theta}_{\text{rev}} = C \mathbf{1}\{\hat{D} \leq -c\} + R \mathbf{1}\{\hat{D} > -c\}.$$

The margin  $c = z_{1-\alpha} s$  with  $s = 2\sigma\|C - R\|/\sqrt{m}$  and  $z_{1-\alpha} = \Phi^{-1}(1-\alpha)$  makes the check a one-sided test at level  $\alpha$  of “the correction does not help” (Lemma 4.1);  $c = 0$  is plain hold-out selection.

We write  $\Delta = \ell(C) - \ell(R)$  for the loss increase caused by the correction (negative when it helps),  $\Delta^+ = \max(\Delta, 0)$ , and  $\pi = \mathbb{P}(\hat{D} \leq -c \mid \mathcal{F}, \theta)$  for the conditional probability that the check passes.

**Shared structure.** In the simulation the parameters of all tasks are exchangeable draws from  $N(\mu, \tau^2 I_d)$  (a hierarchical or empirical Bayes model). The amount of shared structure is measured by  $\text{snr} = \tau^2/\sigma_e^2$ : at  $\text{snr} = 0$  all tasks have the same parameter, at  $\text{snr} \rightarrow \infty$  they share nothing. Departure from the structure is modelled by shifting the target’s centre by  $\kappa\tau$  along one coordinate while the sources stay put.

## 3 What cannot be guaranteed

**Theorem 3.1** (No free lunch for corrections). *Let  $X \sim N(\theta, \sigma_e^2 I_d)$  with  $\theta \in \mathbb{R}^d$  unknown and  $\sigma_e^2$  known, and let  $D$  be auxiliary data, independent of  $X$  given  $\theta$ , whose law  $Q$  does not depend on  $\theta$ . Let  $\delta = \delta(X, D)$  be any estimator;  $D$  may include auxiliary randomisation, so randomised rules are covered, and an estimator with infinite risk satisfies (a) and (c) trivially.*

- (a)  $\sup_\theta \mathcal{R}(\theta, \delta) \geq d\sigma_e^2 = \sup_\theta \mathcal{R}(\theta, X)$ . [classical; proof included]
- (b) If  $\delta = X - b(D)$  for a measurable  $b$  with  $\mathbb{E}\|b(D)\|^2 < \infty$ , then  $\mathcal{R}(\theta, \delta) = d\sigma_e^2 + \mathbb{E}\|b(D)\|^2$  for every  $\theta$ . [proved here]
- (c) If  $d \leq 2$  and  $\mathbb{P}_\theta(\delta \neq X) > 0$  for some (equivalently every)  $\theta$ , then  $\mathcal{R}(\theta', \delta) > d\sigma_e^2$  for some  $\theta'$ . [literature (admissibility) + proved here]
- (d) If  $d \geq 3$  there exist  $\delta$  with  $\mathcal{R}(\theta, \delta) < d\sigma_e^2$  for every  $\theta$ ; by (a) the supremum is still  $d\sigma_e^2$ . [literature]

*Proof.* (a) Fix  $\tau > 0$  and put the prior  $\theta \sim N(0, \tau^2 I_d)$ , so that  $(\theta, X, D)$  has joint law  $\pi(d\theta)P_\theta(dx)Q(dD)$ . Because  $Q$  does not depend on  $\theta$  and  $D$  is independent of  $X$  given  $\theta$ , the conditional law of  $\theta$  given  $(X, D)$  equals its conditional law given  $X$ , namely  $N(\frac{\tau^2}{\tau^2 + \sigma_e^2} X, \frac{\tau^2 \sigma_e^2}{\tau^2 + \sigma_e^2} I_d)$ . The Bayes risk of any  $\delta(X, D)$  is therefore at least that of the posterior mean,  $d\tau^2 \sigma_e^2 / (\tau^2 + \sigma_e^2)$ . Since the supremum of the risk dominates its average,  $\sup_\theta \mathcal{R}(\theta, \delta) \geq d\tau^2 \sigma_e^2 / (\tau^2 + \sigma_e^2)$  for every  $\tau$ ; let  $\tau \rightarrow \infty$ . Equality for  $\delta = X$  is direct. (b)  $\mathbb{E}_\theta \|X - b - \theta\|^2 = \mathbb{E}\|X - \theta\|^2 + \mathbb{E}\|b\|^2 - 2\mathbb{E}\langle X - \theta, b \rangle$ ; the cross term is finite by Cauchy–Schwarz and vanishes because  $X - \theta$  has mean zero and is independent of  $D$ . (c) Let  $\tilde{\delta}(X) = \mathbb{E}[\delta(X, D) | X]$ , which does not depend on  $\theta$  because the conditional law of  $D$  given  $X$  is  $Q$ . Since  $\tilde{\delta}$  is the conditional expectation of  $\delta$  given  $X$  and  $\tilde{\delta} - \theta$  is  $X$ -measurable,  $\mathcal{R}(\theta, \delta) = \mathcal{R}(\theta, \tilde{\delta}) + \mathbb{E}_\theta \|\delta - \tilde{\delta}\|^2$  (Pythagoras; we may assume the left side finite). If  $\mathbb{P}(\tilde{\delta} \neq X) > 0$ , suppose for contradiction that  $\mathcal{R}(\theta, \tilde{\delta}) \leq d\sigma_e^2$  for all  $\theta$ ; then  $\delta' = (\tilde{\delta} + X)/2$  satisfies  $\|\delta' - \theta\|^2 = \frac{1}{2}\|\tilde{\delta} - \theta\|^2 + \frac{1}{2}\|X - \theta\|^2 - \frac{1}{4}\|\tilde{\delta} - X\|^2$ , hence  $\mathcal{R}(\theta, \delta') \leq d\sigma_e^2 - \frac{1}{4}\mathbb{E}_\theta \|\tilde{\delta} - X\|^2 < d\sigma_e^2$  for every  $\theta$  (the Gaussian laws  $P_\theta$  are mutually absolutely continuous, so the last expectation is positive for all  $\theta$  once it is positive for one). This contradicts the admissibility of  $X$  for  $d \leq 2$  [3, 17]; see Lehmann and Casella [12, Ch. 5]. If instead  $\tilde{\delta} = X$  a.s., then  $\mathcal{R}(\theta, \delta) = d\sigma_e^2 + \mathbb{E}_\theta \|\delta - X\|^2 > d\sigma_e^2$  for every  $\theta$ . (d) is the James–Stein theorem [11, 17]; the risk of the James–Stein estimator tends to  $d\sigma_e^2$  as  $\|\theta\| \rightarrow \infty$ .  $\square$

*Remark 3.2* (Where transfer can enter). The only hypothesis of Theorem 3.1(a) that a transfer setting can violate is that the law of the source data does not depend on  $\theta$ . When  $\theta$  and the source parameters are exchangeable draws from  $N(\mu, \tau^2 I_d)$  with  $(\mu, \tau^2)$  fixed, the law of the source data still does not depend on the realised  $\theta$ , so part (a) applies to the frequentist risk even in the hierarchical model; the sources are informative about the task distribution, which is what the Bayes risk over tasks rewards, and that is the natural criterion. For the affine correction  $C = \mu + \lambda(X - \mu)$  with  $\mu$  and  $\lambda \in [0, 1)$  fixed,

$$\mathcal{R}(\theta, C) = \lambda^2 d\sigma_e^2 + (1 - \lambda)^2 \|\theta - \mu\|^2, \quad \mathbb{E}_{\theta \sim N(\mu, \tau^2 I)} \mathcal{R}(\theta, C) = \lambda^2 d\sigma_e^2 + (1 - \lambda)^2 d\tau^2, \quad (1)$$

minimised at  $\lambda^* = \tau^2 / (\tau^2 + \sigma_e^2)$  with value  $d\sigma_e^2 \lambda^* < d\sigma_e^2$ . [classical] The available relative gain over the reference is  $1 - \lambda^* = 1 / (1 + \text{snr})$ . The first formula in (1) is the negative-transfer risk: at a target far from the centre the correction is worse than the reference by  $(1 - \lambda)^2 \|\theta - \mu\|^2 - (1 - \lambda^2) d\sigma_e^2$ , unbounded in  $\theta$ . Reversion is meant to cap this term. Section 4 gives two bounds: one of order  $\mathbb{E}\|C - R\| / \sqrt{m}$ , which for the shrinkage operator grows linearly with  $\|\theta - \mu\|$  and so trades a quadratic harm for a linear one, and a cap uniform in  $\theta$  (Corollary 4.3).

## 4 What can be guaranteed

**Lemma 4.1** (The held-out loss difference). *With the notation of Definition 2.1,  $\hat{D} - \Delta = -2\langle C - R, \bar{Y} - \theta \rangle$ . Hence  $\mathbb{E}[\hat{D} | \mathcal{F}, \theta] = \Delta$ , and if the held-out observations are Gaussian then, conditionally on  $(\mathcal{F}, \theta)$ ,  $\hat{D} \sim N(\Delta, s^2)$  with  $s = 2\sigma\|C - R\| / \sqrt{m}$ . In particular the rule “use  $C$  iff  $\hat{D} \leq -z_{1-\alpha}s$ ” rejects the hypothesis  $\Delta \geq 0$  with conditional probability  $\Phi(-(\Delta + z_{1-\alpha}s)/s) \leq \alpha$  whenever  $\Delta \geq 0$ .*

*Proof.* Expand  $\|C - \bar{Y}\|^2 - \|R - \bar{Y}\|^2 = \|C\|^2 - \|R\|^2 - 2\langle C - R, \bar{Y} \rangle$  and  $\Delta = \|C\|^2 - \|R\|^2 - 2\langle C - R, \theta \rangle$ . Given  $(\mathcal{F}, \theta)$ ,  $\bar{Y} - \theta \sim N(0, \sigma^2 I_d / m)$  is independent of  $C - R$ , so  $\langle C - R, \bar{Y} - \theta \rangle \sim N(0, \sigma^2 \|C - R\|^2 / m)$ .  $\square$

**Theorem 4.2** (Do no harm, up to a controlled amount). *Let  $R, C$  be  $\mathcal{F}$ -measurable, let the held-out sample be independent of  $\mathcal{F}$  given  $\theta$  with mean  $\theta$ , and let  $\hat{\theta}_{\text{rev}}$  be the reverting estimator with margin  $c \geq 0$ . Write  $s = 2\sigma\|C - R\| / \sqrt{m}$ .*

(i) (Identity.)  $\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R) = \mathbb{E}[\Delta \pi] \leq \mathbb{E}[\Delta^+ \pi]$ , where  $\pi = \mathbb{P}(\hat{D} \leq -c | \mathcal{F}, \theta)$ .

(ii) (Gaussian held-out, test at level  $\alpha \in (0, \frac{1}{2}]$ ,  $c = z_{1-\alpha}s$ .) With  $\kappa(\alpha) = \sup_{u \geq 0} u \Phi(-(z_{1-\alpha} + u)) \leq \varphi(z_{1-\alpha})$  (Table 5),

$$\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R) \leq \min \left\{ \alpha \mathbb{E}[\Delta^+], \quad \kappa(\alpha) \frac{2\sigma}{\sqrt{m}} \mathbb{E}\|C - R\| \right\},$$

and the probability of a harmful transfer event ( $C$  is used and  $\Delta > 0$ ) satisfies  $\mathbb{P}(\hat{\theta}_{\text{rev}} = C, \Delta > 0) \leq \alpha \mathbb{P}(\Delta > 0) \leq \alpha$ ; equivalently, given  $\Delta > 0$ ,  $C$  is used with probability at most  $\alpha$ .

(iii) (Bounded held-out.) If instead  $\|Y_j - \theta\| \leq \rho$  a.s. and  $c = 2\rho\|C - R\| \sqrt{2 \log(1/\alpha)/m}$ , then the same two statements hold with  $\kappa(\alpha) 2\sigma / \sqrt{m}$  replaced by  $2\rho / \sqrt{em}$ .

(iv) (Near-oracle selection, Gaussian held-out,  $c = z_{1-\alpha}s$ .)

$$\mathcal{R}(\hat{\theta}_{\text{rev}}) \leq \mathbb{E}[\min\{\ell(R), \ell(C)\}] + (z_{1-\alpha} + \varphi(0)) \frac{2\sigma}{\sqrt{m}} \mathbb{E}\|C - R\|.$$

[proved here; standard arguments]

*Proof.* (i)  $\ell(\hat{\theta}_{\text{rev}}) = \ell(R) + \Delta \mathbf{1}\{\hat{D} \leq -c\}$ . Take expectations conditionally on  $(\mathcal{F}, \theta)$ , on which  $\Delta$  is measurable:  $\mathbb{E}[\ell(\hat{\theta}_{\text{rev}}) - \ell(R) \mid \mathcal{F}, \theta] = \Delta\pi$ . On  $\{\Delta \leq 0\}$  this is  $\leq 0$ , whence the inequality.

(ii) By Lemma 4.1, on  $\{\Delta > 0\}$  we have  $\pi = \Phi(-(z+u))$  with  $z = z_{1-\alpha} \geq 0$  and  $u = \Delta/s \geq 0$  (if  $s = 0$  then  $C = R$ ,  $\Delta = 0$  and there is nothing to prove). First bound:  $\Phi(-(z+u)) \leq \Phi(-z) = \alpha$ , so  $\mathbb{E}[\Delta^+\pi] \leq \alpha\mathbb{E}[\Delta^+]$ . Second bound:  $\Delta^+\pi = s u \Phi(-(z+u)) \leq s \kappa(\alpha)$  by definition of  $\kappa$ ; take expectations. By the Mills ratio inequality  $\Phi(-x) \leq \varphi(x)/x$  for  $x > 0$  and since  $\varphi$  decreases on  $[0, \infty)$ ,

$$u \Phi(-(z+u)) \leq \frac{u}{z+u} \varphi(z+u) \leq \varphi(z+u) \leq \varphi(z),$$

so  $\kappa(\alpha) \leq \varphi(z)$ . For the harmful event,  $\mathbb{P}(\hat{\theta}_{\text{rev}} = C, \Delta > 0) = \mathbb{E}[\mathbf{1}\{\Delta > 0\}\pi] \leq \alpha\mathbb{P}(\Delta > 0)$ .

(iii) Write  $\hat{D} - \Delta = \frac{1}{m} \sum_j W_j$  with  $W_j = -2\langle C - R, Y_j - \theta \rangle$ , independent given  $(\mathcal{F}, \theta)$ , mean zero and  $|W_j| \leq B := 2\rho\|C - R\|$ . Hoeffding's inequality [9] gives  $\mathbb{P}(\hat{D} - \Delta \leq -t \mid \mathcal{F}, \theta) \leq \exp(-mt^2/(2B^2))$ . On  $\{\Delta \geq 0\}$ ,  $\pi \leq \exp(-m(c+\Delta)^2/(2B^2)) \leq \exp(-mc^2/(2B^2)) = \alpha$  for the stated  $c$ , which yields the  $\alpha$ -bound and the bound on the harmful event. Also  $\Delta^+\pi \leq \Delta^+ \exp(-m(\Delta^+)^2/(2B^2)) \leq \sup_{x \geq 0} x e^{-mx^2/(2B^2)} = B/\sqrt{em}$ .

(iv)  $\ell(\hat{\theta}_{\text{rev}}) - \min\{\ell(R), \ell(C)\} = |\Delta| \mathbf{1}\{\text{wrong choice}\}$ , where the choice is wrong on  $\{\Delta > 0, \hat{D} \leq -c\}$  and on  $\{\Delta < 0, \hat{D} > -c\}$ . Conditionally on  $(\mathcal{F}, \theta)$ , with  $u = |\Delta|/s$ : on  $\{\Delta > 0\}$  the conditional expectation is  $s u \Phi(-(z+u)) \leq s \varphi(z+u) \leq s \varphi(0)$  as in (ii); on  $\{\Delta < 0\}$  it is  $s u \Phi(z-u)$  by Lemma 4.1, which is at most  $s z$  when  $u \leq z$  and, writing  $u = z+v$  with  $v > 0$ , at most  $s(z+v\Phi(-v)) \leq s(z+\varphi(v)) \leq s(z+\varphi(0))$  otherwise. Both cases are bounded by  $s(z+\varphi(0))$ ; take expectations. (On  $\{\Delta < 0\}$  the conditional value  $s u \Phi(z-u)$  is at most  $s \sup_{u \geq 0} u \Phi(z-u)$ , which is smaller than  $s(z+\varphi(0))$ ; we keep the closed form.)  $\square$

**Corollary 4.3** (A cap uniform in the parameter). *Under the hypotheses of Theorem 4.2(ii), with  $R = \bar{X}_e$  and any  $\mathcal{F}$ -measurable  $C$ , for every  $\theta$ ,*

$$\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R) \leq \frac{4\sigma}{\sqrt{m}} \kappa(\alpha) \mathbb{E}\|R - \theta\| + \frac{4\sigma^2}{m} \kappa_2(\alpha) \leq 4\sigma^2 \left( \kappa(\alpha) \sqrt{\frac{d}{n_e m}} + \frac{\kappa_2(\alpha)}{m} \right),$$

where  $\kappa_2(\alpha) = \sup_{u \geq 0} u^2 \Phi(-(z_{1-\alpha} + u)) \leq \varphi(1) = \sup_{x \geq 0} x \varphi(x) = 0.2420$ , so that  $(\varphi(z_{1-\alpha}), \varphi(1))$  may replace  $(\kappa, \kappa_2)$  as a closed form. Theorem 4.2(ii) is sharper when  $\mathbb{E}\|C - R\|$  is small; the cap takes over when it is large, where the bound of Theorem 4.2(ii) grows without limit although the harm does not. [proved here]

*Proof.* Write  $a = \|C - R\|$ ,  $e = R - \theta$  and  $z = z_{1-\alpha}$ , so that  $\Delta = \|(C - R) + e\|^2 - \|e\|^2 = a^2 + 2\langle C - R, e \rangle \geq a(a - 2\|e\|)$ ; on  $\{\Delta > 0\}$  put  $u = \Delta/s$ , so that  $\Delta^+\pi = s u \Phi(-(z+u))$ . If  $a \leq 2\|e\|$ , then  $s \leq 4\sigma\|e\|/\sqrt{m}$  and  $\Delta^+\pi \leq s \kappa \leq (4\sigma\|e\|/\sqrt{m})\kappa$ . If  $a > 2\|e\|$ , then  $u \geq w := \sqrt{m}(a - 2\|e\|)/(2\sigma) > 0$  and  $s = (2\sigma/\sqrt{m})a = 4\sigma\|e\|/\sqrt{m} + 4\sigma^2 w/m$ , so, because  $w \leq u$ ,

$$\Delta^+\pi = \frac{4\sigma\|e\|}{\sqrt{m}} u \Phi(-(z+u)) + \frac{4\sigma^2}{m} w u \Phi(-(z+u)) \leq \frac{4\sigma\|e\|}{\sqrt{m}} \kappa + \frac{4\sigma^2}{m} \kappa_2.$$

Take expectations in Theorem 4.2(i) and use  $\mathbb{E}\|e\| \leq (\mathbb{E}\|e\|^2)^{1/2} = \sigma\sqrt{d/n_e}$ . Finally, by the Mills ratio and  $z \geq 0$ ,  $u^2 \Phi(-(z+u)) \leq u \varphi(z+u) \leq (z+u)\varphi(z+u) \leq \varphi(1)$ . Nothing is assumed about  $C$  beyond measurability, and the right side does not depend on  $\theta$ .  $\square$

**Corollary 4.4** (Price of the split). *Let  $R_n = \bar{X}_n$  be the reference computed from all  $n = n_e + m$  target observations, with risk  $d\sigma^2/n$ . Then*

$$\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R_n) = d\sigma^2 \left( \frac{1}{n_e} - \frac{1}{n} \right) + (\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R)),$$

so the total cost of the safety mechanism relative to “do nothing with all the data” is the split cost  $d\sigma^2 m/(nn_e)$ , a fraction  $m/n_e$  of the full-data risk  $d\sigma^2/n$  (equivalently  $m/n$  of the risk of  $R$ ), plus the controlled term of Theorem 4.2. [proved here (arithmetic)]

**Proposition 4.5** (Sharpness). *Fix  $\alpha \in (0, \frac{1}{2}]$  and the check of Theorem 4.2(ii), and let  $u^*$  attain  $\kappa(\alpha)$ .*

- (a) (Best constant.) *For every  $\varepsilon > 0$  there are a Gaussian target model and an  $\mathcal{F}$ -measurable pair  $(R, C)$  with  $\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R) \geq (\kappa(\alpha) - \varepsilon) \frac{2\sigma}{\sqrt{m}} \mathbb{E}\|C - R\|$ . Thus  $\kappa(\alpha)$  is the best constant in Theorem 4.2(ii); the closed form  $\varphi(z_{1-\alpha})$  loses the factor  $\varphi(z_{1-\alpha})/\kappa(\alpha)$ , from 2.3 to 20 in Table 5.*
- (b) (Rate for a fixed estimation sample.) *Let  $R = \bar{X}_e$  and  $c_d = \mathbb{E}\|Z\|$  for  $Z \sim N(0, I_d)$ , so that  $\mathbb{E}\|R - \theta\| = \sigma c_d / \sqrt{n_e}$  and  $c_d \leq \sqrt{d}$ . For all  $d$ ,  $n_e$  and  $m$ , the supremum over  $\theta$  and over  $\mathcal{F}$ -measurable  $C$  of  $\mathcal{R}(\hat{\theta}_{\text{rev}}) - \mathcal{R}(R)$  lies between*

$$4\sigma^2 \left( \frac{\kappa(\alpha) c_d}{\sqrt{n_e m}} + \frac{u^* \kappa(\alpha)}{m} \right) \quad \text{and} \quad 4\sigma^2 \left( \frac{\kappa(\alpha) c_d}{\sqrt{n_e m}} + \frac{\kappa_2(\alpha)}{m} \right).$$

*For a fixed estimation sample the worst-case harm is therefore of exact order  $m^{-1/2}$ , and the first term of Corollary 4.3 is attained; as  $n_e \rightarrow \infty$  it is of order  $1/m$ .*

*Both parts concern this class of checks: a larger margin (in the limit, never using  $C$ ) lowers the worst-case harm and pays through the regret term of Theorem 4.2(iv). The worst case is a correction small relative to the noise of the check ( $\Delta/s = u^* \leq 0.75$ ). [proved here; constants computed numerically, Table 5; (b) also checked by Monte Carlo, Appendix A]*

*Proof.* (a) Take  $R = \bar{X}_e$  and  $C = R + v$  with a fixed vector  $v \neq 0$  (an additive correction fixed in advance, as in Theorem 3.1(b)). Then  $\|C - R\| = \|v\|$ ,  $s = 2\sigma\|v\|/\sqrt{m}$  is deterministic and  $\Delta = \|v\|^2 + 2\langle v, R - \theta \rangle$ . By Theorem 4.2(i) the excess risk equals  $\mathbb{E}[g(\Delta)]$  with  $g(x) = x \Phi(-(x + zs)/s)$ , a continuous function with  $|g(x)| \leq |x|$ . As  $n_e \rightarrow \infty$ ,  $\Delta \rightarrow \|v\|^2$  in  $L^2$ , so  $\mathbb{E}[g(\Delta)] \rightarrow g(\|v\|^2) = s u \Phi(-(z + u))$  with  $u = \|v\|^2/s = \|v\|\sqrt{m}/(2\sigma)$ , which can be given any prescribed value  $u \geq 0$  by the choice of  $\|v\|$ . Choosing  $u = u^*$  and  $n_e$  large gives the claim. (b) The upper bound is Corollary 4.3. For the lower bound fix  $\theta_0$ , let  $e = R - \theta_0$  (nonzero a.s.) and  $C = \theta_0 - (1+t)e$  with  $t = 2\sigma u^*/(\sqrt{m}\|e\|)$ , an  $\mathcal{F}$ -measurable function of  $R$ . At  $\theta = \theta_0$ :  $C - R = -(2+t)e$ ,  $\Delta = ((1+t)^2 - 1)\|e\|^2 = t(2+t)\|e\|^2 > 0$  and  $s = 2\sigma(2+t)\|e\|/\sqrt{m}$ , so  $\Delta/s = t\|e\|\sqrt{m}/(2\sigma) = u^*$  and  $s = 4\sigma\|e\|/\sqrt{m} + 4\sigma^2 u^*/m$ . By Lemma 4.1 and Theorem 4.2(i) the excess risk is  $\mathbb{E}[s u^* \Phi(-(z + u^*))] = \kappa(\alpha)(4\sigma\mathbb{E}\|e\|/\sqrt{m} + 4\sigma^2 u^*/m)$ .  $\square$

**Remark 4.6** (Detectability, not estimability). When the correction helps ( $\Delta < 0$ ), Lemma 4.1 gives the conditional probability of using it,  $\Phi(-z_{1-\alpha} + |\Delta|/s)$  with  $|\Delta|/s = |\Delta|\sqrt{m}/(2\sigma\|C - R\|)$ . In the shift example of Proposition 4.5 with a beneficial  $v$  that cancels the error of  $R$  exactly ( $R - \theta = -v$ , so  $\Delta = -\|v\|^2$ ) this ratio is  $\|v\|\sqrt{m}/(2\sigma)$  whatever the dimension: the check is a one-dimensional test along the direction  $C - R$  supplied by the sources, so a gain can be *verified* on the target with a held-out sample that would be far too small to *learn* the correction there. Gains too small to be detected at level  $\alpha$  are forfeited, up to the regret term of Theorem 4.2(iv).

**Remark 4.7** (Variants of the check). (a) If  $\sigma$  is unknown, the check can be the one-sided  $t$ -test at level  $\alpha$  on the  $m$  projections  $\langle C - R, Y_j \rangle$ ; under Gaussian held-out observations it is exact at  $\Delta = 0$  with power monotone in the non-centrality [13], so the  $\alpha$ -form of Theorem 4.2(ii) and the harmful-event bound hold verbatim. (b) Nothing uses that  $C$  is a function of  $R$ : all statements hold for any two  $\mathcal{F}$ -measurable candidates.

## 5 Relation to known constructs, and the line's conclusion

**Proposition 5.1** (The reverting estimator is penalised hold-out selection). *Let  $\hat{L}_m(\delta) = \frac{1}{m} \sum_{j=1}^m \|\delta - Y_j\|^2$  be the held-out empirical loss. Then  $\hat{L}_m(C) - \hat{L}_m(R) = \hat{D}$ , and*

$$\hat{\theta}_{\text{rev}} = \arg \min_{\delta \in \{R, C\}} \{ \hat{L}_m(\delta) + c \mathbf{1}\{\delta = C\} \}$$

*(ties, a null event, resolved toward  $C$  as in Definition 2.1(c)). With  $c = 0$  it is plain hold-out selection between two candidates. [proved here (one line)]*

*Proof.*  $\hat{L}_m(\delta) = \|\delta - \bar{Y}\|^2 + \frac{1}{m} \sum_j \|Y_j - \bar{Y}\|^2$ , and the second term does not depend on  $\delta$ .  $\square$

Three consequences follow. First, the guarantee of Theorem 4.2 is the two-candidate case of the standard analysis of hold-out selection among finitely many candidates fixed before the held-out sample is

| construct                                  | reverts to                     | trigger                                                | guarantee                                                                                                                                                                                                                |
|--------------------------------------------|--------------------------------|--------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Positive-part James–Stein [2]              | the shrinkage target           | $\ X - \mu\ ^2 < (d - 2)\sigma_e^2$                    | dominates James–Stein; minimax, $d \geq 3$                                                                                                                                                                               |
| Limited-translation rules [5]              | the reference, coordinate-wise | component moved further than a cap                     | bounds the risk increase of each component                                                                                                                                                                               |
| SURE selection [18]                        | either candidate               | sign of an unbiased risk-difference estimate; no split | unbiased estimate; no finite-sample selection bound used here                                                                                                                                                            |
| Source detection in transfer learning [14] | target-only estimator          | data-splitting aggregation on the target               | rates under informative sources; protection against non-informative ones                                                                                                                                                 |
| <b>Reverting estimator (this note)</b>     | the reference $R$              | one-sided test at level $\alpha$ on $\hat{D}$          | Theorem 4.2, Corollary 4.3: harm $\leq \alpha \mathbb{E}\Delta^+$ , $\leq \kappa(\alpha) 2\sigma \mathbb{E}\ C - R\ /\sqrt{m}$ and $\leq$ a cap uniform in $\theta$ ; harmful event $\leq \alpha \mathbb{P}(\Delta > 0)$ |

**Table 1:** The reverting estimator next to constructs it coincides with or resembles. The first two revert in a different direction or by a different mechanism; the third is the same construct with a different check, and plain hold-out selection [1, 4] is the case  $c = 0$  (Proposition 5.1); “negative transfer” is the name given to  $\Delta > 0$  in the transfer-learning literature [16].

seen [4, §8.4; see also Ch. 22], see also Arlot and Celisse [1]; with  $K$  candidates the same argument costs an extra  $\sqrt{\log K}$ , and Theorem 4.2 only sharpens the constants in the Gaussian two-candidate case and makes the one-sided, “do no harm” reading explicit. Second, the correction operator enters the guarantee only through  $\mathbb{E}\|C - R\|$ , the size of the proposed change, and through  $\Delta$ , its actual effect; the guarantee is indifferent to how the operator was obtained. Third, whether the reverting estimator *gains* anything is a question about the shared structure (Remark 3.2) and about whether that gain is detectable with  $m$  held-out points (Remark 4.6). Table 1 places the construction next to the classical ones it resembles.

**What this says about the line’s conclusion.** The public finding of OMR001 was that sufficient differentiation from existing model combinations and operator-based corrections had not been demonstrated. Proposition 5.1 and Theorem 4.2 explain why no differentiation should be expected *from the safety mechanism itself*: safe reversion is held-out selection between two candidates, its only guarantee is a selection guarantee, and that guarantee is available for any pair of candidates. Differentiation, if any, must come from the operator family, which this note does not explore beyond three affine instances, all empirical Bayes rules in the sense of Efron and Morris [6]; whether the line should be closed or continued along the open problems of Section 7 is a decision this note does not take. What it adds is the explicit form of the guarantee with its best constant within the class of fixed-level checks, the uniform cap and its rate, and the finding of the next section about the price of the split.

## 6 Simulation

**Protocol.**  $d = 5$ ,  $\sigma = 1$ ,  $S = 30$  source tasks with  $n_s = 60$  observations each, a target task with  $n = 60$  observations split into  $n_e = 48$  for estimation and  $m = 12$  for the check (a split cost of 25.0% of the full-data risk, 20.0% of the risk of  $R$ ; Corollary 4.4);  $m \in \{3, 6, 12, 24\}$  is also swept, with  $\tau^2$  held at its  $m = 12$  value ( $\tau^2 = 0.0208 = \sigma^2/48$ ) so that only the split moves. Sources and target are exchangeable draws from  $N(\mu, \tau^2 I_d)$  with  $\text{snr} = \tau^2/\sigma_e^2 \in \{0, 0.25, 0.5, 1, 2, 4, 8, 16\}$  (structure sweep), and, at  $\text{snr} = 1$ , the target centre is shifted by  $\kappa\tau$  along one coordinate with  $\kappa \in \{0, 2, 4, 6, 8, 12, 16\}$  (departure sweep). The operators are the empirical Bayes shrinkage  $C = \hat{\mu} + \hat{\lambda}(R - \hat{\mu})$  with  $\hat{\mu}$  the mean of the source means and  $\hat{\lambda} = \hat{\tau}^2/(\hat{\tau}^2 + \sigma_e^2)$ ,  $\hat{\tau}^2 = (\hat{v} - \sigma^2/n_s)^+$ ,  $\hat{v}$  the pooled variance of the source means; full pooling ( $\lambda = 0$ ); and a fixed  $\lambda = \frac{1}{2}$ . The check uses  $\alpha \in \{0.5, 0.2, 0.1, 0.05, 0.01\}$ . Each configuration uses 20 000 replicates with common random numbers across  $\alpha$  and operators, seed 20260930. Two variants *without* a guarantee proved here are also run: *refit*, which applies the same check but then builds the final estimate from all  $n$  observations, so that the independence hypothesis of Theorem 4.2 fails, and *SURE reversion*, which uses no split and applies  $C$  to  $\bar{X}_n$  iff Stein’s unbiased estimate [18] of  $\mathcal{R}(C) - \mathcal{R}(\bar{X}_n)$ , equal to  $(1 - \hat{\lambda}_n)[(1 - \hat{\lambda}_n)\|\bar{X}_n - \hat{\mu}\|^2 - 2d\sigma^2/n]$  for the affine operator, is negative.


**Predefined criteria.** (U) *Useful transfer* at a configuration: the lower end of the paired 95% confidence interval of the relative risk gain with respect to the *full-data* reference  $\bar{X}_n$  is at least 5.0%; the split is charged to the method. (S) *Safety*: the Monte Carlo excess risk of the reverting estimator over  $R$  never

| snr  | $\lambda^*$ | $\bar{X}_n$ | always              | reverting, with guarantee |                     |                     |               | without guarantee   |                     |
|------|-------------|-------------|---------------------|---------------------------|---------------------|---------------------|---------------|---------------------|---------------------|
|      |             |             | all data            | $\alpha=0.5$              | $\alpha=0.2$        | $\alpha=0.1$        | $\alpha=0.01$ | refit 0.1           | SURE                |
| 0    | 0.00        | 0.0833      | 0.0030 <sup>U</sup> | 0.0295 <sup>U</sup>       | 0.0610 <sup>U</sup> | 0.0769 <sup>u</sup> | 0.0996        | 0.0702 <sup>U</sup> | 0.0176 <sup>U</sup> |
| 0.25 | 0.20        | 0.0833      | 0.0233 <sup>U</sup> | 0.0453 <sup>U</sup>       | 0.0697 <sup>U</sup> | 0.0825              | 0.0998        | 0.0743 <sup>U</sup> | 0.0365 <sup>U</sup> |
| 0.5  | 0.33        | 0.0831      | 0.0347 <sup>U</sup> | 0.0542 <sup>U</sup>       | 0.0745 <sup>U</sup> | 0.0853              | 0.1002        | 0.0759 <sup>U</sup> | 0.0449 <sup>U</sup> |
| 1    | 0.50        | 0.0841      | 0.0485 <sup>U</sup> | 0.0660 <sup>U</sup>       | 0.0815              | 0.0897              | 0.1015        | 0.0789 <sup>u</sup> | 0.0562 <sup>U</sup> |
| 2    | 0.67        | 0.0832      | 0.0605 <sup>U</sup> | 0.0773 <sup>u</sup>       | 0.0877              | 0.0933              | 0.1015        | 0.0801              | 0.0653 <sup>U</sup> |
| 4    | 0.80        | 0.0833      | 0.0699 <sup>U</sup> | 0.0868                    | 0.0935              | 0.0973              | 0.1033        | 0.0813              | 0.0728 <sup>U</sup> |
| 8    | 0.89        | 0.0833      | 0.0761 <sup>U</sup> | 0.0933                    | 0.0973              | 0.0997              | 0.1034        | 0.0824              | 0.0776 <sup>u</sup> |
| 16   | 0.94        | 0.0833      | 0.0797              | 0.0976                    | 0.0998              | 0.1014              | 0.1037        | 0.0828              | 0.0804              |

**Table 2:** Structure sweep, EB operator, target drawn from the hierarchy,  $m = 12$ . Risks (mean squared error over  $d = 5$  coordinates; Monte Carlo standard errors are below 0.0005). <sup>U</sup> marks configurations meeting the predefined criterion U (gain  $\geq 5.0\%$  over the full-data reference with 95% confidence); <sup>u</sup>: borderline (lower end within 2 points of the threshold). The reference on the estimation sample has risk  $\approx 0.1042$  against 0.0833 for  $\bar{X}_n$ .

| $\kappa$ | $\bar{X}_n$ | always              | reverting, with guarantee |                     |                     |                     | without guarantee   |                     |
|----------|-------------|---------------------|---------------------------|---------------------|---------------------|---------------------|---------------------|---------------------|
|          |             | all data            | $\alpha=0.5$              | $\alpha=0.2$        | $\alpha=0.1$        | $\alpha=0.01$       | refit 0.1           | SURE                |
| 0        | 0.0837      | 0.0481              | 0.0655                    | 0.0814              | 0.0895 <sup>H</sup> | 0.1012 <sup>H</sup> | 0.0785              | 0.0557              |
| 2        | 0.0842      | 0.0652              | 0.0792                    | 0.0891 <sup>H</sup> | 0.0947 <sup>H</sup> | 0.1025 <sup>H</sup> | 0.0822              | 0.0747              |
| 4        | 0.0831      | 0.1174 <sup>H</sup> | 0.1074 <sup>H</sup>       | 0.1021 <sup>H</sup> | 0.1017 <sup>H</sup> | 0.1031 <sup>H</sup> | 0.0859 <sup>H</sup> | 0.0939 <sup>H</sup> |
| 6        | 0.0837      | 0.2024 <sup>H</sup> | 0.1340 <sup>H</sup>       | 0.1112 <sup>H</sup> | 0.1070 <sup>H</sup> | 0.1048 <sup>H</sup> | 0.0888 <sup>H</sup> | 0.0869 <sup>H</sup> |
| 8        | 0.0834      | 0.3225 <sup>H</sup> | 0.1528 <sup>H</sup>       | 0.1160 <sup>H</sup> | 0.1084 <sup>H</sup> | 0.1043 <sup>H</sup> | 0.0890 <sup>H</sup> | 0.0835 <sup>H</sup> |
| 12       | 0.0838      | 0.6665 <sup>H</sup> | 0.1621 <sup>H</sup>       | 0.1153 <sup>H</sup> | 0.1085 <sup>H</sup> | 0.1050 <sup>H</sup> | 0.0877 <sup>H</sup> | 0.0838              |
| 16       | 0.0835      | 1.1512 <sup>H</sup> | 0.1454 <sup>H</sup>       | 0.1098 <sup>H</sup> | 0.1054 <sup>H</sup> | 0.1042 <sup>H</sup> | 0.0847 <sup>H</sup> | 0.0835              |


**Table 3:** Departure sweep, EB operator,  $\text{snr} = 1$ , target centre shifted by  $\kappa\tau$ ,  $m = 12$ . Risks as in Table 2; Monte Carlo standard errors are below 0.0028. <sup>H</sup>: harmful relative to  $\bar{X}_n$  (upper end of the 95% interval of the gain below zero).



**Figure 1:** Reverting estimators with the guarantee, EB operator, risks relative to the full-data reference. (a) Structure sweep (Table 2): the reverting estimators start from the split cost (dashed) and recover part of the available gain. (b) Departure sweep (Table 3), log scale: always correcting is harmful from  $\kappa = 4$  on; at  $\alpha \leq 0.1$  the reverting estimators stay within 1.07–1.30 times the full-data reference, and plain hold-out ( $\alpha = 0.5$ ) reaches 1.93 at  $\kappa = 12$ . The bound of Theorem 4.2(ii) grows with  $\mathbb{E}\|C - R\|$ , linearly in  $\kappa$  for the shrinkage operator; what is proved for every  $\kappa$  is the cap of Corollary 4.3, which allows up to 1.38 times the full-data reference at  $\alpha = 0.1$  and 2.67 at  $\alpha = 0.5$ . The variants without guarantee are in the last two columns of the tables.

exceeds the identity bound, the  $\alpha$ -bound or the closed form  $\varphi(z)\mathbb{E}[s]$  of Theorem 4.2(i)–(ii) by more than two Monte Carlo standard errors, and the harmful-event frequency never exceeds  $\alpha\hat{\mathbb{P}}(\Delta > 0)$  by more than two binomial standard errors. (H) The frequency of harmful transfer events,  $\hat{\theta}_{\text{rev}} = C$  and  $\Delta > 0$  (the correction is used and its realised loss exceeds that of the reference), is reported for every configuration. Criterion U was fixed in writing before the runs; Appendix A states what was chosen or changed later.

**Results: safety (S).** Across all 345 (configuration, operator,  $\alpha$ ) combinations the bounds of Theorem 4.2(i)–(ii), in the  $\kappa$ -form and in the closed form, hold with the two-standard-error allowance: yes; the harmful-event frequency bound holds: yes; the cap of Corollary 4.3 holds: yes (largest ratio of excess to cap 0.67; at  $m = 12$  and  $\alpha = 0.1$  the cap is 0.0110, 0.11 times the risk of  $R$  and 13.3 times smaller than its closed form 0.1461; the  $\kappa$ -bound exceeds it by up to 2.6 times on the grid). In this design the worst case over operators at  $\alpha = 0.1$  lies between 0.0096 and 0.0106 (Proposition 4.5(b)). Where  $\Delta > 0$  almost surely,



**Figure 2:** Excess risk over  $R$  on the departure sweep ( $\alpha = 0.1$ ) against the bounds of Theorem 4.2, the closed form  $\varphi(z)\mathbb{E}[s]$  and the cap of Corollary 4.3 (dashed). The identity bound  $\mathbb{E}[\Delta^+\pi]$  is nearly attained where the excess is positive; at  $\kappa = 8$  the  $\kappa$ -bound is 0.0067 against an excess of +0.0042; at  $\kappa = 16$  the cap (0.0110) is below the  $\kappa$ -bound (0.0129).

the comparison with  $\mathbb{E}[\Delta^+\pi]$  is a check of identity (i) between two Monte Carlo estimates of the same quantity, not of a bound. The largest ratio of the Monte Carlo excess to the  $\kappa$ -bound is 0.94, to its closed form  $\varphi(z)\mathbb{E}[s]$  0.40, and to the regret bound of Theorem 4.2(iv) it is 0.40. The identity of Theorem 4.2(i) was checked replicate by replicate: paired  $z$ -test where acceptance is frequent (274 combinations, largest  $|z|$  2.83), Poisson 99% interval for the acceptance count where it is rare (71 combinations, 1 failure: full pooling,  $\alpha = 0.05$ ,  $m = 6$ , departure  $8\tau$ , 19 acceptances against 37 expected, not reproduced with 50 fresh seeds, 1882 against 1870,  $z = +0.29$ ); a conditional check with a fixed  $(R, C, \theta)$  agrees with  $\Phi(-(\Delta + zs)/s)$  at every  $\alpha$  (largest  $|z|$  1.72). The frequency of harmful events (the correction used and harmful) of the EB operator never exceeds 25.30% at  $\alpha = 0.5$ , 3.36% at  $\alpha = 0.1$  and 0.25% at  $\alpha = 0.01$ , against up to 100.0% for always correcting on the split sample; the frequency of using  $C$  given that it is harmful never exceeds  $\alpha$  (at most 8.42% for  $\alpha = 0.1$ ). Reversion also caps the two bad operators: full pooling has unreverted risk 1.7324 at  $\text{snr} = 16$  against 0.1064 after reversion (the per-configuration tables for all operators, with every bound, are in `results/tables.md`, T1, T2 and T4).

**Results: usefulness (U).** Always correcting with all the data is useful at  $\text{snr} \in \{0, 0.25, 0.5, 1, 2, 4, 8\}$  and, in the hierarchy, never harmful in the criterion’s sense; its relative gain over  $\bar{X}_n$  is +42.3% at  $\text{snr} = 1$ . The reverting estimator *with* the guarantee is useful at  $\text{snr} \in \{0, 0.25, 0.5, 1, 2\}$  for  $\alpha = 0.5$ ,  $\{0, 0.25, 0.5\}$  for  $\alpha = 0.2$ ,  $\{0\}$  for  $\alpha = 0.1$ , and nowhere for  $\alpha \leq 0.05$ : at  $\text{snr} = 1$  its gain over  $\bar{X}_n$  is +21.5% ( $\alpha = 0.5$ ), -6.6% ( $\alpha = 0.1$ ) and -20.8% ( $\alpha = 0.01$ ), while its gain over the reference on the same estimation sample is +37.1%, +14.5% and +3.1%. The difference is the split cost of Corollary 4.4 (25.0% of the full-data risk, 20.0% of the risk of  $R$ ), which the guarantee does not cover, plus the gain forfeited by the test (Remark 4.6): at  $\text{snr} = 1$  and  $\alpha = 0.1$  the correction is used in 19.6% of the replicates although it helps in 82.3% of them. The held-out-size sweep (`results/tables.md`, T3;  $\tau^2$  fixed, so that the estimators that do not use the split stay flat in  $m$ ) shows the trade-off: with  $m = 3$  (a 5% split) the gain of the guaranteed estimator over  $\bar{X}_n$  is +3.1% at  $\kappa = 0$  and -17.9% at  $\kappa = 8$ , with excess +0.0104 over  $R$  ( $\kappa$ -bound 0.0123); with  $m = 24$  it is -25.0% and -68.8%. The two variants without guarantee are useful over a wider range: refit at  $\text{snr} \in \{0, 0.25, 0.5, 1\}$  ( $\alpha = 0.1$ ) and SURE reversion at  $\text{snr} \in \{0, 0.25, 0.5, 1, 2, 4, 8\}$ , the latter within a few points of always correcting in the hierarchy. 4 of the marks in Table 2 are borderline (lower end of the interval within 2 points of the threshold):  $\alpha = 0.1$  at  $\text{snr} = 0$  (gain +7.6%, lower end +6.9%); refit 0.1 at  $\text{snr} = 1$  (gain +6.1%, lower end +5.8%);  $\alpha = 0.5$  at  $\text{snr} = 2$  (gain +7.1%, lower end +6.5%); SURE at  $\text{snr} = 8$  (gain +6.9%, lower end +6.6%); in particular “refit useful up to  $\text{snr} = 1$ ” should not be read as robust (the reduced run, `-fast` with 4000 replicates, keeps every other mark and drops that one).

**Results: negative transfer.** On the departure sweep always correcting with all the data becomes harmful (upper end of the 95% interval of the gain below zero) at  $\kappa \in \{4, 6, 8, 12, 16\}$  and reaches 13.79 times the full-data reference risk at  $\kappa = 16$ . The guaranteed reverting estimator at  $\alpha = 0.1$  stays between 1.07 and 1.30 times the full-data reference over the whole sweep; its excess over  $R$  is +0.0042 at  $\kappa = 8$  against the bounds 0.0044 (identity), 0.0067 ( $\kappa$ ) and 0.0296 ( $\alpha$ ), and +0.0012 at  $\kappa = 16$  against 0.0129 ( $\kappa$ ) and the cap 0.0110. The unguaranteed variants are also capped in this design (SURE reversion is harmful only at  $\kappa \in \{4, 6, 8\}$ , by at most 1.13 times the reference), but nothing proved here says they must be.

**Verdict under the predefined criterion.** Safety holds everywhere, as it must. Useful transfer with the guaranteed construction occurs only when the tasks are close ( $\text{snr} \leq 2$  with plain hold-out, only  $\text{snr} = 0$  at  $\alpha = 0.1$ , never at  $\alpha \leq 0.05$ ), because elsewhere the split costs more than the transfer gains.

The constructions useful over a wide range are classical (empirical Bayes shrinkage with SURE selection; cross-validated selection and refit), with no finite-sample “do no harm” bound proved here.

## 7 Claims and their status

| claim                                                                                                                                                                                                                                                                                                                                                       | status                                                                                       |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------|
| No estimator using source data whose law is free of $\theta$ beats the minimax risk $d\sigma_e^2$ (a); a fixed additive correction is worse by exactly $\mathbb{E}\ b\ ^2$ at every $\theta$ (b); for $d \leq 2$ any rule that ever corrects is worse than $R$ somewhere (c); for $d \geq 3$ pointwise but not worst-case improvement exists (d) (Thm. 3.1) | (a) classical, proof included; (b) proved here; (c) literature + proved here; (d) literature |
| Bayes and frequentist risk of affine shrinkage; available gain $1/(1 + \text{snr})$ (Rem. 3.2)                                                                                                                                                                                                                                                              | classical                                                                                    |
| Held-out loss difference conditionally unbiased, Gaussian with sd $2\sigma\ C - R\ /\sqrt{m}$ (Lem. 4.1); excess-risk identity $\mathbb{E}[\Delta\pi]$ (Thm. 4.2(i))                                                                                                                                                                                        | proved here; identity checked in 345 combinations                                            |
| Harm $\leq \min\{\alpha\mathbb{E}\Delta^+, \kappa(\alpha)2\sigma\mathbb{E}\ C - R\ /\sqrt{m}\}$ , closed form $\varphi(z) \geq \kappa(\alpha)$ ; $\mathbb{P}(C \text{ used}, \Delta > 0) \leq \alpha\mathbb{P}(\Delta > 0)$ ; Hoeffding version; near-oracle selection within $(z + \varphi(0))2\sigma\mathbb{E}\ C - R\ /\sqrt{m}$ (Thm. 4.2(ii)–(iv))     | proved here; every bound holds everywhere (S)                                                |
| Cap uniform in $\theta$ with constants $(\kappa, \kappa_2)$ , for any measurable $C$ (Cor. 4.3)                                                                                                                                                                                                                                                             | proved here; holds everywhere                                                                |
| Total cost relative to the full-data reference = split cost ( $m/n_e$ of its risk) + controlled term (Cor. 4.4)                                                                                                                                                                                                                                             | proved here                                                                                  |
| Within one-sided checks at a fixed level $\alpha$ : $\kappa(\alpha)$ is the best constant; for a fixed estimation sample the worst-case harm over operators is of exact order $m^{-1/2}$ (Prop. 4.5)                                                                                                                                                        | proved here; constants numerical; (b) checked by Monte Carlo                                 |
| Reverting estimator = penalised hold-out selection between two candidates; hence no differentiation is to be expected from the safety mechanism itself (Prop. 5.1, Sec. 5)                                                                                                                                                                                  | proved here                                                                                  |
| Differentiation through the operator family                                                                                                                                                                                                                                                                                                                 | not explored (three affine instances)                                                        |
| Useful transfer (criterion U) only at $\text{snr} \in \{0, 0.25, 0.5, 1, 2\}$ ( $\alpha = 0.5$ ), $\{0, 0.25, 0.5\}$ (0.2), $\{0\}$ (0.1), never for $\alpha \leq 0.05$ , in this design                                                                                                                                                                    | verified, design-specific; 4 marks borderline                                                |
| Refit and SURE reversion useful over a wider range and capped in this design                                                                                                                                                                                                                                                                                | empirical; no guarantee                                                                      |
| General accuracy improvement, transfer guarantee, statements about real populations                                                                                                                                                                                                                                                                         | not claimed                                                                                  |

**Table 4:** Claims and their status. “Verified” and “holds” mean reproduced by the script of Appendix A with the fixed seed; neither is a proof.

**Limitations.** The model is Gaussian location with known variance; the  $\kappa$ -bound and the cap use Gaussian held-out observations, the Hoeffding version needs bounded ones, and the  $t$ -test variant is only remarked on. The reference is the sample mean; for a biased reference (a shared instrument bias, say) the held-out sample of the target carries the same bias and the check cannot see the correction, so that setting needs unbiased calibration data and is not covered. The simulation has one design ( $d = 5$ ,  $n = 60$ ,  $S = 30$ ) and one structural family; the “useful” region depends on the split fraction and on the dimension through the ratio of the  $\kappa$ -bound to the reference risk, of order  $\sqrt{n_e/m}/\sqrt{d}$  for the shrinkage operator, and would move in another design. The unguaranteed variants are reported, not analysed. Nothing is said about real data.

**Next steps for the research line.** (a) A finite-sample bound for the refit variant, in which the held-out sample enters both the check and the final estimate, or for a cross-fitted version; this is where the split cost could be recovered with a guarantee. (b) A selection bound for SURE reversion with the affine operator, whose selection region is a ball around  $\hat{\mu}$  and whose risk is not given by SURE because of the discontinuity on the sphere. (c) If the original operator definitions of OMR001 are recovered, place each in Table 1 and compute  $\mathbb{E}\|C - R\|$  for it; by Theorem 4.2 that is all the safety analysis needs, and such a family is the only remaining source of differentiation. (d) Decide whether the line should be closed with this note or continued only along (a)–(b).

## A Reproducibility

The folder contains `experiments/safe_reversion.py` (all sweeps, constants, identity re-checks, tables and figures; `-fast` runs a reduced version) and `experiments/make_numbers.py`, which converts `results/results.json` into the macros and table bodies of this document, so that every number in the text is generated and none is typed. The full run used seed 20260930 (re-checks: 20260930 + 1 and 20260930 + 1000 +  $i$ ), Python 3.11.15, NumPy 2.4.6 [8], SciPy 1.17.1 [19] and Matplotlib 3.11.2 [10], and took 14 s

of wall-clock time (13 s of CPU) on a shared four-core machine. `experiments/check_uniform_cap.py` (seconds; no number in the text comes from it) checks Corollary 4.3 on a grid and Proposition 4.5(b) by Monte Carlo.

**Procedure.** Criterion U was fixed in writing before the runs (script `docstring` and `results.json`). The two-standard-error tolerance of S and the thresholds of the identity check ( $|z| \leq 3$ , Poisson 99%) were chosen before the final run but not recorded beforehand. The Poisson comparison replaced a paired  $z$ -test after a first run showed that the normal approximation fails where acceptance is rare. After the first internal review the held-out-size sweep was re-run with  $\tau^2$  fixed and the borderline mark was introduced; after the second, checks against the  $\kappa$ -form of Theorem 4.2(ii) and against the cap with  $(\kappa, \kappa_2)$  were added, and maxima quoted as “never exceeds” are rounded up; no simulated number changed. No other change to the protocol or the criteria was made after seeing results. Proposition 4.5(b) and the constants  $\kappa, \kappa_2$  in Theorem 4.2(ii) and Corollary 4.3 are due to the second internal review. Risks are Monte Carlo means with standard errors from the replicate standard deviation; relative gains use paired differences and treat the denominator as fixed (its relative standard error is below 1%);  $\kappa(\alpha)$  and  $\kappa_2(\alpha)$  are computed by `scipy.optimize.minimize_scalar` on  $[0, 20]$ , which contains both maximisers (at most 0.75 and 1.19), with tolerance  $10^{-10}$ . The draws per replicate are listed in the README; the constants in Table 5.

| $\alpha$ | $z_{1-\alpha}$ | $\varphi(z_{1-\alpha})$ | $\kappa(\alpha)$ | $u^*$ | $\kappa/\varphi(z_{1-\alpha})$ | $\kappa_2(\alpha)$ | $\kappa_2/\varphi(1)$ | $z_{1-\alpha} + \varphi(0)$ |
|----------|----------------|-------------------------|------------------|-------|--------------------------------|--------------------|-----------------------|-----------------------------|
| 0.5      | 0.000          | 0.3989                  | 0.1700           | 0.752 | 0.426                          | 0.1657             | 0.685                 | 0.399                       |
| 0.2      | 0.842          | 0.2800                  | 0.0451           | 0.543 | 0.161                          | 0.0331             | 0.137                 | 1.241                       |
| 0.1      | 1.282          | 0.1755                  | 0.0188           | 0.465 | 0.107                          | 0.0120             | 0.049                 | 1.680                       |
| 0.05     | 1.645          | 0.1031                  | 0.0082           | 0.413 | 0.079                          | 0.0047             | 0.019                 | 2.044                       |
| 0.01     | 2.326          | 0.0267                  | 0.0013           | 0.337 | 0.049                          | 0.0006             | 0.003                 | 2.725                       |

**Table 5:** Constants of Theorem 4.2, Corollary 4.3 and Proposition 4.5:  $\kappa(\alpha) = \sup_{u \geq 0} u\Phi(-(z_{1-\alpha} + u))$ , attained at  $u^*$ , and  $\kappa_2(\alpha) = \sup_{u \geq 0} u^2\Phi(-(z_{1-\alpha} + u))$ ;  $\varphi(1) = 0.2420$ .

## References

- [1] Sylvain Arlot and Alain Celisse. A survey of cross-validation procedures for model selection. *Statistics Surveys*, 4: 40–79, 2010.
- [2] A. J. Baranchik. A family of minimax estimators of the mean of a multivariate normal distribution. *The Annals of Mathematical Statistics*, 41(2):642–645, 1970.
- [3] Colin R. Blyth. On minimax statistical decision procedures and their admissibility. *The Annals of Mathematical Statistics*, 22(1):22–42, 1951.
- [4] Luc Devroye, László Györfi, and Gábor Lugosi. *A Probabilistic Theory of Pattern Recognition*. Springer, New York, 1996.
- [5] Bradley Efron and Carl Morris. Limiting the risk of Bayes and empirical Bayes estimators—Part II: The empirical Bayes case. *Journal of the American Statistical Association*, 67(337):130–139, 1972.
- [6] Bradley Efron and Carl Morris. Stein’s estimation rule and its competitors—an empirical Bayes approach. *Journal of the American Statistical Association*, 68(341):117–130, 1973.
- [7] Bradley Efron and Carl Morris. Data analysis using Stein’s estimator and its generalizations. *Journal of the American Statistical Association*, 70(350):311–319, 1975.
- [8] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, et al. Array programming with NumPy. *Nature*, 585: 357–362, 2020.
- [9] Wassily Hoeffding. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association*, 58(301):13–30, 1963.
- [10] John D. Hunter. Matplotlib: A 2D graphics environment. *Computing in Science & Engineering*, 9(3):90–95, 2007.
- [11] Willard James and Charles Stein. Estimation with quadratic loss. In *Proceedings of the Fourth Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pages 361–379. University of California Press, Berkeley, 1961.
- [12] Erich L. Lehmann and George Casella. *Theory of Point Estimation*. Springer, New York, 2nd edition, 1998.
- [13] Erich L. Lehmann and Joseph P. Romano. *Testing Statistical Hypotheses*. Springer, New York, 3rd edition, 2005.
- [14] Sai Li, T. Tony Cai, and Hongzhe Li. Transfer learning for high-dimensional linear regression: Prediction, estimation and minimax optimality. *Journal of the Royal Statistical Society: Series B*, 84(1):149–173, 2022.
- [15] Carl N. Morris. Parametric empirical Bayes inference: Theory and applications. *Journal of the American Statistical Association*, 78(381):47–55, 1983.
- [16] Sinno Jialin Pan and Qiang Yang. A survey on transfer learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10):1345–1359, 2010.
- [17] Charles Stein. Inadmissibility of the usual estimator for the mean of a multivariate normal distribution. In *Proceedings of the Third Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pages 197–206. University of California Press, Berkeley, 1956.
- [18] Charles M. Stein. Estimation of the mean of a multivariate normal distribution. *The Annals of Statistics*, 9(6): 1135–1151, 1981.
- [19] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, et al. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nature Methods*, 17:261–272, 2020.