

Fragility of Lasso Variable Selection:

Removal Witnesses, Exact Deletion Tests, Non-Monotonicity and Computational Status

Luciano Silva Alarco

Pontificia Universidad Católica del Perú

ORCID 0000-0002-3395-3129

Working draft v0.2 (one round of internal review applied) — research line TCD001 — October 3, 2026

Abstract

We study how the set of variables selected by the Lasso changes when observations are removed. For data D , a penalty μ and a rule fixing the penalty on subsamples, a *removal witness* for a target change (a variable leaves the selected set, one enters, or the set changes at all) is a set R of observations such that the Lasso on $D \setminus R$ exhibits the change; the *fragility number* is the minimal $|R|$. We establish: (i) an exact closed-form test of whether removing R preserves the signed support (KKT conditions and a Woodbury update), which for $|R| = 1$ under a constant penalty is a Cook–Belsley deletion formula with the Lasso residual, with a cheap certificate against single removals and a polynomial but loose one for k removals; (ii) non-monotonicity of witness families, by hand for $p = 1$ and in exact arithmetic for $p = 2$, where the mechanism is competition between variables; (iii) NP-completeness (weak) of the existence of a deselection witness already for $p = 1$, while the selection direction is solvable in $O(n \log n)$; (iv) on two synthetic Gaussian families with $n \leq 14$, exact fragility numbers by exhaustive search (the selected set changes after one removal in 92% and 52% of the instances) and the failure of three heuristics under a criterion fixed in advance, because at $n \leq 14$ none is cheaper than one tenth of the exhaustive search. No novelty is claimed for the algebra.

1 Introduction

Variable selection with the Lasso [19] is known to be unstable: small changes of the data can change the selected set. This instability has been addressed statistically, by aggregating selections over random subsamples [stability selection, 14], and asymptotically, by conditions under which the selected set is the true one with high probability [5, 21, 23]. A different, finite-sample and combinatorial question is the subject of the research line to which this draft belongs (TCD001): *how many observations must be removed, and which, for the selected set to change?* The question has a long history for the coefficients of least squares, where case-deletion diagnostics [3, 6] and the influence curve [11] measure the effect of dropping one observation, and where the worst-case effect of dropping a small fraction of the data has recently been studied under the name of approximate maximum influence perturbation [4], through adaptive heuristics for influential sets [13], and through exact or certified algorithms for the sign of a regression coefficient [8, 15]. For a *selection*, the quantity of interest is discrete, so the influence function does not apply directly, and the natural object is a set of observations, a *witness*, together with its minimal size.

The public record of the line states that non-monotone changes and connections to known combinatorial problems were identified, and that novelty, computational advantage and some certificates were not sufficiently established. This draft gives the line a precise core: it proves what can be proved, reports which certificates are loose, and shows that on small synthetic instances the predefined advantage criterion is not met because exhaustive search with the closed-form test is already cheap; novelty is not claimed. Section 2 fixes the definitions (penalty rules on subsamples, removal witnesses, fragility numbers, monotonicity). Section 3 proves the exact deletion test (Proposition 3.1), its single-removal form, a Cook–Belsley deletion formula with the Lasso residual (Corollary 3.2), and a polynomial-time certificate that no k removals change the support (Proposition 3.4), all verified numerically against refits. Section 4 shows that witness families are not closed under supersets, with a hand-checkable example for $p = 1$ and an exactly verified example for $p = 2$ driven by competition between variables. Section 5 gives the computational status: existence of a deselection witness is NP-complete for $p = 1$ by a reduction from SUBSET SUM, selection witnesses for $p = 1$ are easy, and the case $p \geq 2$ is open. Section 6 reports exact fragility numbers by exhaustive search, witness types, non-monotone pairs, scaling with n , and three heuristics against the exact minimum under a criterion fixed before the runs. The paper is methodological: it makes no claim about real data and does not propose

the fragility number as a replacement for stability selection (Remark 3.5 states the one implication that follows from the definitions); every number in the text is produced by the scripts of Appendix A with fixed seeds, and Section 7 lists all claims with their status.

2 Setting and definitions

Let $D = (X, y)$ with $X \in \mathbb{R}^{n \times p}$, $y \in \mathbb{R}^n$, no intercept and columns as given. For $\mu > 0$ the Lasso objective in sum form is

$$L_D(\beta; \mu) = \frac{1}{2} \|y - X\beta\|_2^2 + \mu \|\beta\|_1, \quad (1)$$

which is the objective of scikit-learn’s **Lasso** with **alpha** = μ/n [17]. We write $\hat{\beta}(D; \mu)$ for its minimiser when it is unique, $S(D) = \text{supp } \hat{\beta}$ for the *support* (the selected set) and (S, s) with $s = \text{sgn } \hat{\beta}_S \in \{\pm 1\}^S$ for the *signed support*. Two facts from Tibshirani [20] are used throughout: if the columns of X are in general position the minimiser is unique for every y and $\mu > 0$; and if the *equicorrelation set* $E = \{j : |X_j^\top (y - X\hat{\beta})| = \mu\}$ has X_E of full column rank, the minimiser is unique and is determined by the linear system of the KKT conditions on E . Conversely a unique minimiser has X_S of full column rank (Lemma 2.4). [literature]

Definition 2.1 (Penalty rule). A *penalty rule* assigns to the subsample $D \setminus R$ of size $n - k$ a penalty μ_k . Two rules are used: the *constant rule* (C), $\mu_k = \mu$, and the *proportional rule* (P), $\mu_k = \mu(n - k)/n$, which keeps scikit-learn’s **alpha** fixed and is what a practitioner does who refits “with the same λ ” in the mean-squared-error parametrisation.

Definition 2.2 (Removal witness, fragility number). Fix D , μ , a rule, and let (S, s) be the signed support on D . A *target* Π is a property of the signed support of the Lasso on a subsample, relative to (S, s) ; the targets used here are

- ANY: $S(D \setminus R) \neq S$ (unsigned), or $(S(D \setminus R), s(D \setminus R)) \neq (S, s)$ (signed);
- LEAVE(j), $j \in S$: $j \notin S(D \setminus R)$;
- ENTER(j), $j \notin S$: $j \in S(D \setminus R)$.

A *removal witness* for Π is a nonempty proper subset $\emptyset \neq R \subsetneq [n]$ such that the Lasso on $D \setminus R$ with penalty $\mu_{|R|}$ has a unique minimiser and Π holds. The *witness family* is $W_\Pi(D) = \{R : R \text{ is a witness}\}$ and the *fragility number* is $f_\Pi(D) = \min\{|R| : R \in W_\Pi(D)\} \geq 1$, with $f_\Pi(D) = \infty$ if $W_\Pi(D) = \emptyset$.

The signed and unsigned versions of ANY differ only when a coefficient changes sign without leaving the support; $f_{\text{ANY}}^{\text{signed}} \leq f_{\text{ANY}}$ and the experiments report how often the two differ. Witnesses for the signed target are decided without refitting by Proposition 3.1.

Definition 2.3 (Monotonicity). The target Π is *monotone on* D if $R \in W_\Pi(D)$ and $R \subseteq R' \subsetneq [n]$ imply $R' \in W_\Pi(D)$. A *non-monotone pair* is a pair $R \subsetneq R'$ with $R \in W_\Pi(D)$ and $R' \notin W_\Pi(D)$; for LEAVE(j) this means that j leaves the support after removing R and re-enters after removing R' .

Lemma 2.4. [classical] *If the minimiser of (1) is unique with support S , then X_S has full column rank.*

Proof. If $X_S \eta = 0$ with $\eta \neq 0$, extend η by zeros. For $|t|$ small, $\hat{\beta} + t\eta$ has the same fit and the same signs on S , so $L_D(\hat{\beta} + t\eta) = L_D(\hat{\beta}) + t\mu s^\top \eta$. If $s^\top \eta \neq 0$ one sign of t decreases the objective, contradicting optimality; if $s^\top \eta = 0$ every such $\hat{\beta} + t\eta$ is a minimiser, contradicting uniqueness. (This is part of the characterisation in 20.) $\square$

3 An exact deletion test and two certificates

Throughout this section $\hat{\beta}$ is the unique minimiser on D with signed support (S, s) , $|S| = s_0 \geq 0$ (we write s_0 for the size to avoid a clash with the sign vector), X_S has full column rank, and

$$\begin{aligned} A &= X_S^\top X_S, & r &= y - X\hat{\beta}, & c &= X^\top r, & V &= X_S A^{-1} \text{ (rows } v_i^\top), \\ h_i &= x_{iS}^\top A^{-1} x_{iS}, & \tilde{X} &= (I - X_S A^{-1} X_S^\top) X_S^c, \end{aligned}$$

so that $c_S = \mu s$, $|c_j| \leq \mu$ for $j \notin S$, $h_i \in [0, 1]$ are the leverages of the active-set regression and $\tilde{X}_j$ is the residual of X_j regressed on X_S . For $R \subseteq [n]$ with $|R| = k$ let $U = X_{R,S} \in \mathbb{R}^{k \times s_0}$ and $H_R = U A^{-1} U^\top$, the $R \times R$ block of the projection $X_S A^{-1} X_S^\top$.

Proposition 3.1 (Exact removal test). [proved here] Let $R \subsetneq [n]$, $|R| = k$, and let μ' be the penalty on $D \setminus R$, $\delta = \mu - \mu'$.

(a) If $I - H_R$ is singular, (S, s) is not the signed support of a unique minimiser on $D \setminus R$.

(b) If $I - H_R$ is nonsingular, put $e_R = (I - H_R)^{-1}r_R$, $w = A^{-1}s$, $z = w + V_R^\top(I - H_R)^{-1}Uw$, and

$$\tilde{\beta}_S = \hat{\beta}_S - V_R^\top e_R + \delta z, \quad (2)$$

$$\tilde{c}_j = c_j - \tilde{X}_{R,j}^\top e_R - \delta(X_j^\top X_S - X_{R,j}^\top U)z, \quad j \notin S. \quad (3)$$

Then (S, s) is the signed support of a minimiser on $D \setminus R$ if and only if $s_j \tilde{\beta}_j > 0$ for all $j \in S$ and $|\tilde{c}_j| \leq \mu'$ for all $j \notin S$. If moreover $|\tilde{c}_j| < \mu'$ for all $j \notin S$, the minimiser on $D \setminus R$ is unique and equals $\tilde{\beta}$ (extended by zeros).

Under the constant rule $\delta = 0$ and the last terms of (2)–(3) vanish.

Proof. A vector β with signed support (S, s) minimises the objective on $D \setminus R$ if and only if it satisfies the KKT conditions

$$X_{S,-R}^\top(y_{-R} - X_{S,-R}\beta_S) = \mu' s, \quad |X_{j,-R}^\top(y_{-R} - X_{S,-R}\beta_S)| \leq \mu' \quad (j \notin S), \quad s_j \beta_j > 0 \quad (j \in S). \quad (4)$$

The matrix of the linear system is $A' = X_{S,-R}^\top X_{S,-R} = A - U^\top U$, and by the matrix determinant lemma $\det A' = \det A \cdot \det(I - UA^{-1}U^\top) = \det A \cdot \det(I - H_R)$.

(a) If $I - H_R$ is singular then $X_{S,-R}$ is rank deficient, and by Lemma 2.4 no unique minimiser on $D \setminus R$ has support S .

(b) If $I - H_R$ is nonsingular, the Sherman–Morrison–Woodbury identity [10, 18, 22] gives $A'^{-1} = A^{-1} + A^{-1}U^\top(I - H_R)^{-1}UA^{-1}$, and the unique solution of the linear system is $\beta_S = A'^{-1}(X_S^\top y - U^\top y_R - \mu' s)$. Write $X_S^\top y - \mu s = A\hat{\beta}_S$ (the KKT equation on D) so that $X_S^\top y - U^\top y_R - \mu' s = A\hat{\beta}_S - U^\top y_R + \delta s$. Then

$$\begin{aligned} A'^{-1}(A\hat{\beta}_S - U^\top y_R) &= \hat{\beta}_S + A^{-1}U^\top[(I - H_R)^{-1}U\hat{\beta}_S - y_R - (I - H_R)^{-1}H_R y_R] \\ &= \hat{\beta}_S - A^{-1}U^\top(I - H_R)^{-1}(y_R - U\hat{\beta}_S), \end{aligned}$$

using $y_R + (I - H_R)^{-1}H_R y_R = (I - H_R)^{-1}y_R$; since $y_R - U\hat{\beta}_S = r_R$ and $A^{-1}U^\top = V_R^\top$, this is $\hat{\beta}_S - V_R^\top e_R$. The remaining term $\delta A'^{-1}s = \delta z$ by the same identity, which proves (2). For $j \notin S$, with $\beta_S = \hat{\beta}_S - V_R^\top e_R$ the full-length residual is $\tilde{r} = r + X_S V_R^\top e_R$, whose R -block is $r_R + H_R e_R = (I - H_R)^{-1}r_R = e_R$; hence

$$\begin{aligned} X_{j,-R}^\top \tilde{r} - R &= X_j^\top \tilde{r} - X_{R,j}^\top e_R = c_j + X_j^\top X_S A^{-1}U^\top e_R - X_{R,j}^\top e_R \\ &= c_j - (X_{R,j} - UA^{-1}X_S^\top X_j)^\top e_R = c_j - \tilde{X}_{R,j}^\top e_R, \end{aligned}$$

and the additional coefficient δz contributes $-\delta X_{j,-R}^\top X_{S,-R} z = -\delta(X_j^\top X_S - X_{R,j}^\top U)z$, which proves (3). The equivalence is now (4). If the inactive inequalities are strict, the equicorrelation set on $D \setminus R$ is S and $X_{S,-R}$ has full column rank, so the minimiser is unique [20]. $\square$

The test costs $O(k^3 + pk)$ per set R after the $O(nps_0 + s_0^3)$ precomputation of V , h , $\tilde{X}$ and c ; it is vectorised over all subsets of a given size in the code, and Section 6 measures its throughput against refits. Boundary cases ($s_j \tilde{\beta}_j = 0$ or $|\tilde{c}_j| = \mu'$) have probability zero for continuous designs and are counted, not hidden, in the validation.

The test decides whether (S, s) is preserved, which is not quite the same as deciding whether R is a witness. For the signed ANY target, failure of the test (or case (a)) identifies a witness in the sense of Definition 2.2 provided the minimiser on $D \setminus R$ is unique: for designs with a continuous distribution this holds for every R with probability one (the columns of X_{-R} are in general position, 20), which is how the exhaustive searches of Section 6 use the test; for integer data (Example 4.3) uniqueness is certified separately in exact arithmetic. For the other targets the flagged sets are refitted and the target is read from the new support.

Corollary 3.2 (Single removal; a stability margin). [proved here] Under the constant rule let $R = \{i\}$ and suppose $h_i < 1$ (equivalently, $X_{S,-i}$ has full column rank; $h_i = 1$ is the rank-deficient case of Proposition 3.1(a), in which (S, s) is not preserved). Put $e_i = r_i/(1 - h_i)$. Then

$$\tilde{\beta}_S = \hat{\beta}_S - v_i e_i, \quad \tilde{c}_j = c_j - \tilde{x}_{ij} e_i \quad (j \notin S),$$

so removing observation i preserves (S, s) if and only if $s_j(\hat{\beta}_j - v_{ij}e_i) > 0$ for $j \in S$ and $|c_j - \tilde{x}_{ij}e_i| < \mu$ for $j \notin S$ (boundary cases aside). With $m = \min_{j \in S} |\hat{\beta}_j|$ and $\gamma = \min_{j \notin S} (\mu - |c_j|)$ (if $S^c = \emptyset$, $\gamma = +\infty$ and the second term below is 0; $\gamma = 0$ is a KKT boundary case and is excluded), define the instability index

$$\iota_i = |e_i| \max \left\{ \frac{\|v_i\|_\infty}{m}, \frac{\max_{j \notin S} |\tilde{x}_{ij}|}{\gamma} \right\} \quad \text{if } h_i < 1, \quad \iota_i = +\infty \quad \text{if } h_i = 1.$$

Then $\iota_i < 1$ implies that removing i preserves (S, s) , and $\max_i \iota_i < 1$ implies $f_{\text{ANY}}^{\text{signed}}(D) \geq 2$. Both statements are checkable in $O(np)$ time after the precomputation.

Proof. Specialise Proposition 3.1 to $k = 1$: $H_R = h_i$ and $I - H_R = 1 - h_i$ is nonsingular exactly when $h_i < 1$, since $\det A' = (1 - h_i) \det A$. The sufficient condition is the triangle inequality applied to (2)–(3); when $h_i = 1$ the removal does not preserve (S, s) by part (a), consistently with $\iota_i = +\infty$. The hypothesis is not vacuous: for $n = p = 2$, $X = I_2$, $y = (3, 3)^\top$, $\mu = 1$ the minimiser $\hat{\beta} = (2, 2)^\top$ is unique and $h_1 = h_2 = 1$. $\square$

Remark 3.3 (Relation to case-deletion diagnostics). Here $r = y - X\hat{\beta}$ is the residual of the Lasso fit on D , which differs from the least-squares residual of y on X_S by $\mu X_S A^{-1} s$, and $e_i = r_i / (1 - h_i)$ is the leave-one-out residual of the fixed-support fit. Under the constant rule, $v_i e_i = A^{-1} x_{iS} e_i$ is the DFBETA formula of Belsley et al. [3] for the regression of y on X_S with r_i in place of the least-squares residual (under rule P the extra term δz appears); $\tilde{c}_j - c_j = -\tilde{x}_{ij} e_i$ has no least-squares counterpart and is what decides entries. Cook’s distance [6] of the active-set regression is proportional to $h_i e_i^2$ and is used below as a target-agnostic heuristic. [classical]

Proposition 3.4 (Polynomial certificate for k removals, constant rule). [proved here] For a vector $a \geq 0$ let $T_k(a)$ be the sum of its k largest entries. Suppose $T_k(h) < 1$. Then for every R with $|R| \leq k$ the matrix $I - H_R$ is nonsingular and, with $\tilde{\beta}$ as in (2) and $\delta = 0$,

$$\begin{aligned} |\tilde{\beta}_j - \hat{\beta}_j| &\leq B_j(k) := T_k(|v_{\cdot j}| \odot |r|) + \sqrt{T_k(v_{\cdot j}^2) T_k(r^2)} \frac{T_k(h)}{1 - T_k(h)} \quad (j \in S), \\ |\tilde{c}_j - c_j| &\leq C_j(k) := T_k(|\tilde{x}_{\cdot j}| \odot |r|) + \sqrt{T_k(\tilde{x}_{\cdot j}^2) T_k(r^2)} \frac{T_k(h)}{1 - T_k(h)} \quad (j \notin S). \end{aligned}$$

Consequently, if $B_j(k) < |\hat{\beta}_j|$ for all $j \in S$ and $C_j(k) < \mu - |c_j|$ for all $j \notin S$, then no removal of at most k observations changes the signed support: $f_{\text{ANY}}^{\text{signed}}(D) \geq k + 1$. The certificate costs $O(nps_0 + pn \log n)$ (the T_k of $p + 1$ vectors of length n are read off after sorting each once; $O(np)$ with selection instead of sorting) and is monotone in k ; we write $k^*(D)$ for the largest certified k ($k^* = 0$ if none).

Proof. H_R is a principal submatrix of a projection, so $0 \preceq H_R \preceq I$ and $\lambda_{\max}(H_R) \leq \text{tr } H_R = \sum_{i \in R} h_i \leq T_k(h) < 1$; hence $I - H_R$ is nonsingular with $\|(I - H_R)^{-1}\|_2 \leq (1 - T_k(h))^{-1}$ and $\|e_R\|_2 \leq \|r_R\|_2 / (1 - T_k(h)) \leq \sqrt{T_k(r^2)} / (1 - T_k(h))$. From $(I - H_R)e_R = r_R$ we have $e_R = r_R + H_R e_R$, so

$$\begin{aligned} \tilde{\beta}_j - \hat{\beta}_j &= -v_{R,j}^\top e_R = -\sum_{i \in R} v_{ij} r_i - v_{R,j}^\top H_R e_R, \\ |v_{R,j}^\top H_R e_R| &\leq \|v_{R,j}\|_2 \|H_R\|_2 \|e_R\|_2 \leq \frac{\sqrt{T_k(v_{\cdot j}^2) T_k(r^2)} T_k(h)}{1 - T_k(h)}, \end{aligned}$$

and the first term is bounded by $T_k(|v_{\cdot j}| \odot |r|)$. The bound for $\tilde{c}_j$ is identical with $\tilde{x}_{\cdot j}$ in place of $v_{\cdot j}$. Under the hypotheses, every R with $|R| \leq k$ satisfies the conditions of Proposition 3.1(b) strictly, so the signed support is preserved. $\square$

The first term of $B_j(k)$ is the additive, “sum of single-deletion effects” approximation on which sorted-influence heuristics rest [4]; the second term controls the interaction between removed observations. Since $\sum_i h_i = s_0$, the hypothesis $T_k(h) < 1$ restricts the certificate to $k < n/s_0$ roughly, and the experiments show that it is loose: it certifies $k^* \geq 1$ in only a fraction of the instances whose true fragility number is at least 2 (Section 6).

Remark 3.5 (Stability selection). [proved here] Stability selection [14] refits on subsamples of size $\lceil n/2 \rceil$ with a fixed penalty and reports the selection frequency Π_j of variable j ; whether the fixed penalty is rule C or rule P depends on the parametrisation (sum form, as in the original paper, or per-observation form, as in common software). Under whichever rule is used, a subsample of size $\lceil n/2 \rceil$ is $D \setminus R$ with $|R| = \lceil n/2 \rceil$, so $f_{\text{LEAVE}(j)}(D) > \lceil n/2 \rceil$ under that rule implies $\Pi_j = 1$. The converse cannot be obtained from the definitions: $f_{\text{LEAVE}(j)} \leq \lceil n/2 \rceil$ produces a witness R of that size or smaller, but by Section 4 a superset of R of size exactly $\lceil n/2 \rceil$ need not be a witness. We did not search for an instance in which the converse actually fails.

data set	support	exact coefficients	exact inactive c_j	marginal $x_1^\top y, x_2^\top y$
D	$\{1\}$	$\hat{\beta}_1 = \frac{3}{20}$	$c_2 = -\frac{1}{10}$	5, -1
$D \setminus R, R = \{3, 4\}$	$\{2\}$	$\hat{\beta}_2 = -\frac{5}{12}$	$c_1 = \frac{11}{4}$	4, -6
$D \setminus R', R' = \{2, 3, 4\}$	$\{1\}$	$\hat{\beta}_1 = \frac{1}{10}$	$c_2 = -\frac{17}{10}$	4, -2

Table 1: Example 4.3 in exact arithmetic ($\mu = \frac{7}{2}$, rule C). Coefficients solve the KKT system on the listed support; inactive correlations $c_j = X_j^\top (y - X\hat{\beta})$ satisfy $|c_j| < \mu$ strictly, which certifies the unique minimiser.

4 Non-monotonicity

Proposition 4.1. *For each of the targets LEAVE, ENTER and ANY there exist data D and μ for which the target is not monotone, under both penalty rules.*

Example 4.2 ($p = 1$, by hand). [proved here] Let $x = (1, 1, 1)^\top$, $y = (2, -2, 2)^\top$, $\mu = \frac{3}{2}$. For $p = 1$ the minimiser is $\hat{\beta} = \text{sgn}(x^\top y)(|x^\top y| - \mu)_+ / \|x\|^2$, so the variable is selected iff $|x^\top y| > \mu$. On D , $x^\top y = 2 > \frac{3}{2}$: selected. On $D \setminus \{1\}$, $x^\top y = 0$: not selected under rule C ($\mu' = \frac{3}{2}$) and under rule P ($\mu' = 1$). On $D \setminus \{1, 2\}$, $x^\top y = 2$: selected under rule C ($2 > \frac{3}{2}$) and under rule P ($\mu'' = \frac{1}{2}$). Thus $\{1\}$ is a LEAVE witness and its superset $\{1, 2\}$ is not: $R = \{1\}$, $R' = \{1, 2\}$ is a non-monotone pair for LEAVE and for ANY. For ENTER take $x = (1, 1, 1, 1)^\top$, $y = (2, -2, 2, -2)^\top$, $\mu = \frac{3}{2}$: the variable is not selected on D ($x^\top y = 0$), is selected after removing observation 2 ($x^\top y = 2$, above $\mu' = \frac{3}{2}$ under rule C and $\mu' = \frac{9}{8}$ under rule P), and is not selected after removing $\{2, 3\}$ ($x^\top y = 0$); both removal sets are proper. Both examples are checked in exact arithmetic by the script of Appendix A (12/12 data-set/rule combinations).

The one-variable example only says that a sum can leave an interval and come back. The next example shows the mechanism specific to selection: a variable is excluded not because its own correlation with the response drops below the penalty, but because a competing variable absorbs it.

Example 4.3 ($p = 2$, competition; exact arithmetic). [computationally verified (exact arithmetic)] Let $n = 5$, $p = 2$, $\mu = \frac{7}{2}$ and

$$[X \mid y] = \begin{bmatrix} 2 & -1 & 2 \\ 0 & 2 & -2 \\ 1 & 1 & 3 \\ 2 & -2 & -1 \\ -1 & 1 & 0 \end{bmatrix} \quad (\text{rows are observations; last column is } y),$$

with rule C, $R = \{3, 4\}$ and $R' = \{2, 3, 4\}$. Table 1 lists, for D , $D \setminus R$ and $D \setminus R'$, the signed support returned by the homotopy solver and the exact rational solution of the KKT system on that support, with all inequalities checked strictly in exact arithmetic (so the support is that of the unique minimiser). Variable 1 is selected on D , leaves on $D \setminus R$ while variable 2 enters, and returns on $D \setminus R'$. The exclusion is due to competition: on $D \setminus R$ the marginal correlation $|x_1^\top y| = 4$ exceeds $\mu = \frac{7}{2}$, so with $p = 1$ variable 1 would be selected, but the residual correlation after fitting variable 2 is $c_1 = \frac{11}{4}$, below μ . The instance was found by a seeded search over integer data in $\{-2, \dots, 2\}^{5 \times 2} \times \{-3, \dots, 3\}^5$ and rational μ (11 trials); all 26 proper removal sets of this instance are exactly verifiable (no KKT ties), and among them there are 3 non-monotone pairs for LEAVE(1).

Non-monotonicity means that a minimal witness cannot be found by growing a set until the target holds and then pruning; exhaustive search by increasing size, or a heuristic followed by verification, are the honest options at small n , and a bound on the fragility number says nothing about subsample frequencies beyond Remark 3.5.

5 Computational status

For $p = 1$, with $a_i = x_i y_i$ and $T = \sum_i a_i$, the variable is selected on $D \setminus R$ iff $|\sum_{i \notin R} a_i| > \mu_{|R|}$. The problem is therefore about the sums of subfamilies of the numbers a_i , and the two directions of change behave very differently.

Theorem 5.1 ($p = 1$). [proved here] *Consider instances with $p = 1$, rational data and rational $\mu > 0$, under either penalty rule.*

- (a) **DESELECTION-WITNESS**: given (x, y, μ) with the variable selected on D , decide whether some proper $R \subsetneq [n]$ deselects it. This problem is NP-complete. Consequently computing f_{LEAVE} (equivalently the unsigned f_{ANY} , since for $p = 1$ every unsigned change is a deselection) is NP-hard, already for $p = 1$, hence for the general problem.
- (b) **SELECTION-WITNESS**: given (x, y, μ) with the variable not selected on D , the minimal $|R|$ such that it is selected on $D \setminus R$ can be computed in $O(n \log n)$ time.

Proof. (a) Membership in NP: given R , compute $\sum_{i \notin R} a_i$ and compare with $\mu_{|R|}$. Hardness: reduce from SUBSET SUM [9, 12]: given positive integers $b_1, \dots, b_m$ and a target $t \geq 1$, decide whether some subfamily sums to t ; without loss of generality $t < \sum_i b_i$ (otherwise the answer is immediate). Set $n = m + 1$, $x_i = 1$ for all i , $y_i = b_i$ for $i \leq m$, $y_{m+1} = -t$ and $\mu = \frac{1}{2}$. Then $a = (b_1, \dots, b_m, -t)$ and $T = \sum_i b_i - t \geq 1 > \mu \geq \mu_k$ for every k , so the variable is selected on D . For a proper R with $K = [n] \setminus R \neq \emptyset$, the variable is deselected on $D \setminus R$ iff $|\sum_{i \in K} a_i| \leq \mu_{|R|} \leq \frac{1}{2}$, iff $\sum_{i \in K} a_i = 0$ because the sum is an integer. If $m + 1 \notin K$ the sum is a nonempty sum of positive integers; hence $m + 1 \in K$ and $\sum_{i \in K \setminus \{m+1\}} b_i = t$, which is a solution of the SUBSET SUM instance, and conversely every solution K' gives the witness $R = [n] \setminus (K' \cup \{m + 1\})$, proper because $t < \sum_i b_i$ forces $K' \neq [m]$. The reduction is polynomial and works verbatim for rule P, since $\mu_k = \frac{1}{2}(n - k)/n < 1$.

(b) Suppose $|T| \leq \mu$. For k fixed, over all R with $|R| = k$ the sum $\sum_{i \notin R} a_i$ ranges over $[T - \sigma_k^+, T - \sigma_k^-]$ where σ_k^+ and σ_k^- are the sums of the k largest and of the k smallest a_i , and both endpoints are attained. Hence a witness of size k exists iff $T - \sigma_k^- > \mu_k$ or $T - \sigma_k^+ < -\mu_k$. After sorting the a_i once, all $\sigma_k^\pm$ are prefix sums and the smallest feasible $k \leq n - 1$ is found in $O(n)$. $\square$

Remark 5.2 (What the theorem does and does not say). (i) The hardness is inherited from SUBSET SUM and is therefore *weak*: for integer data a dynamic programme over partial sums computes f_{LEAVE} in time polynomial in n and in $\sum_i |a_i|$ (with the cardinality as a second state under rule P). The difficult instances are those in which the deselection window $[-\mu_k, \mu_k]$ is narrow relative to the granularity of the data; strong NP-hardness (hardness for numbers polynomially bounded in n) is excluded for $p = 1$ by this pseudo-polynomial algorithm unless $P = NP$. (ii) For $p \geq 2$ the theorem gives NP-hardness of the general LEAVE and ANY problems only because $p = 1$ instances are instances of the general problem; we have no hardness proof for ENTER with $p \geq 2$ and no polynomial algorithm for any target with $p \geq 2$. (iii) The analogous question for the sign of a least-squares coefficient is the subject of Moitra and Rohatgi [15] and Freund and Hopkins [8], and maximum feasible subsystems of linear relations are treated by Amaldi and Kann [1, 2]; we rely on no hardness statement from those papers and do not claim that the reduction above is new. [literature / not claimed]

Conjecture 5.3. [conjectural] For $p \geq 2$, SELECTION-WITNESS (existence of a proper R such that a given inactive variable enters $S(D \setminus R)$) is NP-complete.

6 Experiments

Protocol. All experiments use Gaussian designs with rows $x_i \sim N(0, \Sigma)$, $\Sigma_{jk} = \rho^{|j-k|}$, β^* with two nonzero entries $\pm b$, noise $N(0, 1)$, and $\mu = c\sqrt{2n \log p}$ (the universal threshold in sum form, scaled by c). Two families are used (their choice is discussed under “Pre-specification and changes” below): **A** (“strong signal”), $p = 4$, $b = 4$, $\rho = 0.0$, $c = 1.5$; **B** (“noisy”), $p = 8$, $b = 2$, $\rho = 0.5$, $c = 1.0$. The Lasso is solved by the homotopy algorithm [7, 16] as implemented in `sklearn.linear_model.lars_path`, whose output satisfies the KKT equalities to 10^{-14} and agrees in support with coordinate descent at tolerance 10^{-13} on all 104 refits checked (maximal coefficient difference 2×10^{-13}). Rule C is the default; rule P is reported where stated. Seeds are 20260930 (E0–E3, E5) and 20260931 (E4). Exhaustive searches enumerate removal sets by increasing size, decide the signed target with Proposition 3.1 and refit only the flagged sets when the target requires the new support; targeted searches are capped at $|R| \leq 5$ and the ANY search at $|R| \leq n - 2$ ($|R| \leq 6$ in E4), and a capped search that finds nothing is reported as “not found”, never as a value.

Criterion for “computational advantage” (fixed before the runs). A heuristic has computational advantage on a family if (i) it returns a witness of exactly minimal size in at least 90% of the instances whose exact fragility number is finite within the cap, and (ii) its median wall-clock cost is at most 10% of that of the exhaustive search on the same instances and machine.

Pre-specification and changes. The criterion above was fixed in the code (`run_fragility.py`) before any stored run, is recorded in `results/fragility.json`, and was never moved. Two choices were made after looking at data: (i) Families A and B were chosen after a pilot sweep, not included in the results, over the penalty factor c at $n \leq 14$, in which almost all instances with $c \leq 1$ had $f = 1$; A is the most stable family found whose support is not always the true one, B a noisy one. The fragility fractions reported below therefore describe these two families and nothing beyond them. (ii) After a first complete run of E2, E3 and E5 the one-step greedy was corrected (its score became the predicted margin of the target itself instead of any break of the current pattern, which had favoured untargeted changes) and E2, E3 and E5 were rerun; E1 and E4 do not use the modified branch. For this version all four scripts were rerun with the final library in a single session, so that `results/` comes from one version of the code. *Timing.* The exhaustive-search timer starts after the initial fit and the $O(nps_0)$ precomputation, whereas the heuristic timers include the initial fit (about one refit, 442 μ s at $n = 14$); removing one refit's worth of time from every heuristic leaves the smallest median cost ratio over all combinations at 0.26 (raw: 0.46), so no verdict changes.

6.1 E1: validation of Proposition 3.1 and of the certificates

On four (n, p) cells ($b = 2$, $\rho = 0.3$, $c = 1$), with 15 instances per cell for the comparison with refits and 30 fresh instances per cell for the certificates, Proposition 3.1 was compared with a refit for every single removal and for random sets of 2, 3 and 5 observations, under both rules. Agreement was 2760/2760 for single removals and 7200/7200 for sets, and whenever the signed support was preserved the closed-form coefficients (2) matched the refit to 4×10^{-15} (single) and 6×10^{-15} (sets); the per-cell counts are in `results/validation.md`. The sufficient condition of Corollary 3.2 certified 1216 of the 2114 single removals that are in fact stable (58%), and never a removal that is not. On $n = 14$, $p = 6$, all 1001 subsets of size 4 were tested at 1.1 μ s per subset against 442 μ s per refit, a factor of 399. The per-cell results for the certificate of Proposition 3.4 (fraction of instances with $f \geq 2$ in which $k^* \geq 1$, mean gap $f - 1 - k^*$) are in `results/validation.md`; they are summarised with E2 below.

6.2 E2: exact fragility numbers for $n \leq 14$

For 60 instances per cell we computed the exact signed fragility number for ANY under both rules, the type of change realised by the minimal witnesses, and the targeted fragility numbers for the weakest selected variable (smallest $|\hat{\beta}_j|$) and for the strongest inactive variable (largest $|c_j|$); the per-cell table is `results/fragility.md`, pooled values follow. In family B the selected set changes after a single removal in 92% of the 240 instances (maximal $f = 3$); in family A in 52% (maximal $f = 4$). Family B had 8 instances with an empty support, 2 of them without a witness within $|R| \leq n - 2$ (counted as not found); signed and unsigned ANY differed in 0 (A) and 0 (B) instances. Among the minimal witnesses examined, the change is a single variable leaving in 51% (A) and 54% (B) of the cases, a single variable entering in 43% and 25%, and a simultaneous leave and enter in 6% and 20%; sign-only changes occurred 0 and 0 times. Targeted searches were not concluded within the cap in 12/240 (A, LEAVE) and 31/237 (A, ENTER) instances. The certificate of Proposition 3.4 gave $k^* \geq 1$ in 30% of the 115 family-A instances with $f \geq 2$ (mean gap 1.2).

6.3 E3: heuristics against the exact minimum

Three heuristics were run on the same instances (rule C): *one-step greedy*, which at every step removes the observation whose exact single-removal effect (Corollary 3.2 on the current fit) brings the target closest, refits, and repeats; *sorted scores*, in the style of the approximate maximum influence perturbation [4], which scores every observation once on the full data by its exact single-removal effect on each margin, predicts the number of removals k_{pred} from cumulative sums, and then removes observations in that fixed order with a refit after each until the target holds; and *Cook's distance*, target-agnostic. Table 2 pools the results over $n \leq 14$ against the predefined criterion. 0 of the 18 (family, target, heuristic) combinations met it. The one-step greedy is the most accurate: it returns a minimal witness in 97%/100% of the instances for ANY (families A/B), 95%/98% for LEAVE and 86%/88% for ENTER (lower on the instances with $f \geq 2$, Table 2; maximal excess over the minimum in family A: 4). These are point estimates on 206–240 instances with standard errors of 0.4–3.1 points: the one-step greedy meets the accuracy part of the criterion for ANY and LEAVE, while for ENTER its shortfall (86%/88% against 90%) is 1.7 and 1.1 standard errors below the threshold and should not be read as a settled failure. The sorted-scores heuristic is less accurate (82%/100%, 75%/95%, 71%/85%) and Cook's distance is the worst for the targeted changes. No heuristic meets the cost part: at these sizes none is cheaper than one tenth of the exhaustive search that uses Proposition 3.1 (smallest median cost ratio over all combinations 0.46, or 0.26 after removing the heuristics' initial fit), with median cost ratios of the one-step greedy of 1.34/0.73 (ANY), 0.48/0.58 (LEAVE) and 0.76/0.52 (ENTER).

fam.	target	heuristic	n	exact $\pm$ s.e.	n ($f \geq 2$)	exact ($f \geq 2$)	cost ratio	adj. cost ratio	advantage
A	any	one-step greedy	240	97% $\pm$ 1.2	115	93%	1.34	0.81	no
A	any	sorted scores (AMIP-style)	240	82% $\pm$ 2.5	115	63%	1.28	0.77	no
A	any	Cook's distance	240	66% $\pm$ 3.1	115	57%	1.23	0.77	no
A	leave	one-step greedy	228	95% $\pm$ 1.4	167	93%	0.48	0.31	no
A	leave	sorted scores (AMIP-style)	228	75% $\pm$ 2.9	167	65%	0.46	0.31	no
A	leave	Cook's distance	228	67% $\pm$ 3.1	167	66%	0.49	0.31	no
A	enter	one-step greedy	206	86% $\pm$ 2.4	133	78%	0.76	0.48	no
A	enter	sorted scores (AMIP-style)	206	71% $\pm$ 3.1	133	56%	0.72	0.45	no
A	enter	Cook's distance	206	27% $\pm$ 3.1	133	11%	1.04	0.75	no
B	any	one-step greedy	238	100% $\pm$ 0.0	17	100%	0.73	0.37	no
B	any	sorted scores (AMIP-style)	238	100% $\pm$ 0.4	17	94%	0.77	0.40	no
B	any	Cook's distance	238	87% $\pm$ 2.2	17	76%	0.62	0.33	no
B	leave	one-step greedy	232	98% $\pm$ 0.9	48	96%	0.58	0.29	no
B	leave	sorted scores (AMIP-style)	232	95% $\pm$ 1.5	48	79%	0.56	0.28	no
B	leave	Cook's distance	232	78% $\pm$ 2.7	48	65%	0.53	0.29	no
B	enter	one-step greedy	209	88% $\pm$ 2.3	69	80%	0.52	0.28	no
B	enter	sorted scores (AMIP-style)	209	85% $\pm$ 2.5	69	72%	0.49	0.26	no
B	enter	Cook's distance	209	23% $\pm$ 2.9	69	1%	0.78	0.56	no

Table 2: E3. Heuristics against the exact minimum, pooled over $n \in \{8, 10, 12, 14\}$, rule C. “Exact” is the fraction of instances (with finite exact f within the cap) in which the heuristic’s witness has minimal size; “s.e.” is the binomial standard error in percentage points; “cost ratio” is the median of heuristic time over exhaustive-search time, and “adj.” removes one refit (the heuristic’s initial fit) from the heuristic time. The verdict applies the criterion fixed before the run (exact $\geq 90\%$ and median cost ratio $\leq 10\%$).

The linear prediction k_{pred} of the sorted-scores heuristic equals the exact fragility number in 70%/99% of the ANY instances (A/B) and underestimates it in 0%/0%.

6.4 E4: scaling with n

With the closed-form test the exhaustive search for signed ANY extends to $n = 34$ at $|R| \leq 6$ (Figure 1a,b; per-cell table in `results/scaling.md`). In family A the censored mean fragility number (instances above the cap excluded, which biases the mean downwards as n grows; medians are in `results/scaling.md`) rises slowly with n (1.43 at $n = 10$, 4.00 at $n = 34$, where 1 of 20 instances lie above the cap), while in family B it stays close to one (75% of the instances have $f = 1$ at $n = 34$). Over the 233 family-A instances with finite f the one-step greedy returned the exact minimum in 88% of the cases (83% of the 174 with $f \geq 2$, maximal excess 6). At $n = 34$ its median cost fell to 0.070 of the exhaustive search (0.41 s on average) while it was exact in 89% of the 20 instances: the heuristic becomes cheap at the sizes where its accuracy drops below the criterion (over the 113 family-A instances with $n \geq 22$ and finite f the median cost ratio is 0.774, or 0.68 over all 120 instances with $n \geq 22$, and the exactness 84%; ratios at $n \leq 18$, where the exhaustive search takes under 10 ms, are dominated by timer resolution and are not quoted). Beyond exhaustive reach, Figure 1c,d brackets the fragility number between the certificate of Proposition 3.4 ($k^* + 1$) and the greedy upper bound g : at $n = 400$ in family A the median greedy witness has 10 observations (mean $g/n \approx 0.028$), the median certified lower bound is $k^* + 1$ with $k^* = 5$, and the mean ratio $(k^* + 1)/g$ is 0.57; the bracket is wide, and the observed decrease of g/n with n is an empirical statement about these families only.

6.5 E5: frequency of non-monotone pairs

For every instance of E2 with a finite targeted fragility number, all one-element supersets $R \cup \{i\}$ of up to 20 minimal witnesses R were refitted. For LEAVE of the weakest variable, the variable had re-entered the support in 7.1% of the 6960 superset pairs of family A (at least one re-entry in 51% of the 228 instances) and in 10.5% of 8092 pairs of family B (72% of 232 instances). For ANY, the original support was restored by adding one observation to a minimal witness in 14.9% (A) and 9.3% (B) of the pairs (per-cell values in `results/fragility.md`). Non-monotonicity is thus not a curiosity of constructed examples but a regular feature of minimal witnesses on these designs.

7 What can be claimed

Limitations. The exact test presupposes a unique minimiser and full column rank on the support; ties are measure-zero for continuous designs but real in integer data, and the code counts them rather than resolving them. For the signed ANY target the exhaustive search records a non-preserved set as a witness without refitting, which presupposes uniqueness of the minimiser on $D \setminus R$; this holds almost surely for the Gaussian




Figure 1: (a,b) Distribution of the exact signed fragility number for ANY (rule C) against n , families A and B, from E2 ($n \leq 14$) and E4 ($n \geq 18$; search capped at $|R| \leq 6$, the top segment collects $f \geq 6$ and not found). (c,d) Bracketing at large n (rule C): median one-step greedy witness size (upper bound, with interquartile range), median certified lower bound $k^* + 1$ from Proposition 3.4, and exact medians from E4 where available.

claim	status
A unique Lasso minimiser has X_S of full column rank; uniqueness under full rank of X_E (Lemma 2.4)	classical / literature
Exact closed-form test of whether removing R preserves the signed support, both rules (Prop. 3.1); agreement with refits in 2760/2760 single and 7200/7200 multiple removals (E1); identifies signed-ANY witnesses when the minimiser on $D \setminus R$ is unique (a.s. for continuous designs)	proved here; verified
Single removal with $h_i < 1$ (rule C) is a DFBETA-type formula with the Lasso residual; $O(np)$ certificate for $f \geq 2$ (Cor. 3.2)	proved here
Polynomial certificate for $f \geq k + 1$ (Prop. 3.4); certifies $k^* \geq 1$ in 30% of family-A instances with $f \geq 2$, $n \leq 14$: loose (E1, E2, E4)	proved here; verified
Witness families are not monotone: hand-checkable example with $p = 1$ (Ex. 4.2); competition example with $p = 2$, exact rational KKT certificate (Ex. 4.3)	proved here; verified (exact arithmetic)
$f_{\text{LEAVE}(j)} > \lceil n/2 \rceil$ implies stability-selection frequency 1; converse not implied (Rem. 3.5)	proved here
Existence of a deselection witness is NP-complete for $p = 1$, weakly; fragility number NP-hard in general; selection witnesses for $p = 1$ in $O(n \log n)$ (Thm. 5.1)	proved here
NP-completeness of selection witnesses for $p \geq 2$ (Conj. 5.3)	conjectural
Single-removal fragility ($f = 1$) in 92% (B) and 52% (A) of instances with $n \leq 14$; minimal witnesses mostly single events (E2); non-monotone pairs among one-element supersets of minimal witnesses in a measurable fraction of cases (E5)	verified, these two families only
No heuristic meets the predefined advantage criterion at $n \leq 14$ (all fail the cost part; the one-step greedy's accuracy shortfall for ENTER is within two s.e.); the one-step greedy is the most accurate and becomes cheap only for $n \geq 22$ (E3, E4)	verified, these two families only
Rates of f/n with n ; behaviour on real data; novelty of the algebra or of the reduction	not claimed

Table 3: Claims and their status. “Verified” means reproduced by the scripts of Appendix A with the fixed seeds; it is not a proof and it is specific to the two synthetic families.

designs used here and is certified exactly for the integer examples, but not in general for integer data. The homotopy solver mishandles exact ties of correlations (two variables entering at the same penalty, as in the instance of the proof of Corollary 3.2); every fit is therefore checked against the KKT conditions with a coordinate-descent fallback, which the reference runs never triggered. The two families were chosen after a pilot sweep (Section 6), so the fragility fractions are statements about these families only. The certificate of Proposition 3.4 is loose and restricted to $k < n/s_0$. The hardness statement is for $p = 1$ and is weak in the

technical sense; it says nothing about the difficulty of typical instances, on which exhaustive search with the closed-form test was fast. The experiments use two synthetic families, one solver, and a cap on the size of removal sets; the fragility numbers reported are exact under that cap and are not estimates of anything beyond the families. The comparison with heuristics is on a single machine and the cost ratios depend on the implementation (a vectorised test against a Python-level loop of refits).

Next steps for the research line. (a) Prove or refute Conjecture 5.3, and settle the complexity of LEAVE for fixed $p \geq 2$ in the strong sense (the $p = 1$ argument only gives weak hardness). (b) Tighten Proposition 3.4 by exploiting the sign structure of the interaction term or by a small semidefinite relaxation, and measure the gap to f on E4. (c) Replace the fixed- μ setting by a data-driven rule (cross-validation), for which the witness must account for the change of μ itself. (d) Extend the exact test to the elastic net and to the square-root Lasso, whose KKT systems are also linear on a fixed signed support.

A Reproducibility

The library `experiments/lasso_fragility.py` and the scripts `exact_examples.py` (E0), `run_validation.py` (E1), `run_fragility.py` (E2, E3, E5), `run_scaling.py` (E4), `make_figures.py` and `make_numbers.py` (which turns `results/*.json` into the macros of this document, so that no number in the text is typed) are described in the README. The reference run used Python 3.11.15, NumPy 2.4.6, SciPy 1.17.1, scikit-learn 1.9.1 [17], seeds 20260930 and 20260931, and took 7 s (E1), 40 s (E2/E3/E5) and 34 s (E4) on one core of a shared four-core machine, 1.3 minutes in all.

Solver. `lars_path(X, y, method="lasso", alpha_min= $\mu/n'$ )` on the reduced data with n' rows; the last column of the path is the minimiser of (1) at penalty μ . Supports are read off exact zeros of the homotopy output (threshold 10^{-10}). Every fit is KKT-checked (relative tolerance 10^{-8}) and recomputed by coordinate descent if the check fails, which happens only on exact ties of correlations (`check_b1_counterexample.py`); the fallback never fired in the reference runs.

Exact test and search. For a batch of subsets of equal size k : H_R by one batched matrix product, e_R by a batched solve, $\det(I - H_R) \leq 10^{-12}$ treated as rank deficient (a change), conditions of Proposition 3.1(b) with strict inequalities. Sizes $k = 1, 2, \dots$ in order; for the signed ANY target no refit is needed, for the other targets flagged subsets are refitted; up to 50 minimal witnesses are collected. Heuristics and exact-arithmetic checks are documented in the docstrings.

References

- [1] Edoardo Amaldi and Viggo Kann. The complexity and approximability of finding maximum feasible subsystems of linear relations. *Theoretical Computer Science*, 147(1–2):181–210, 1995.
- [2] Edoardo Amaldi and Viggo Kann. On the approximability of minimizing nonzero variables or unsatisfied relations in linear systems. *Theoretical Computer Science*, 209(1–2):237–260, 1998.
- [3] David A. Belsley, Edwin Kuh, and Roy E. Welsch. *Regression Diagnostics: Identifying Influential Data and Sources of Collinearity*. Wiley, New York, 1980.
- [4] Tamara Broderick, Ryan Giordano, and Rachael Meager. An automatic finite-sample robustness metric: when can dropping a little data make a big difference?, 2020. arXiv:2011.14999, v4 (2023).
- [5] Peter Bühlmann and Sara van de Geer. *Statistics for High-Dimensional Data: Methods, Theory and Applications*. Springer, Berlin, 2011.
- [6] R. Dennis Cook. Detection of influential observation in linear regression. *Technometrics*, 19(1):15–18, 1977.
- [7] Bradley Efron, Trevor Hastie, Iain Johnstone, and Robert Tibshirani. Least angle regression. *The Annals of Statistics*, 32(2):407–499, 2004.
- [8] Daniel Freund and Samuel B. Hopkins. Towards practical robustness auditing for linear regression, 2023. arXiv:2307.16315.
- [9] Michael R. Garey and David S. Johnson. *Computers and Intractability: A Guide to the Theory of NP-Completeness*. W. H. Freeman, San Francisco, 1979.
- [10] William W. Hager. Updating the inverse of a matrix. *SIAM Review*, 31(2):221–239, 1989.
- [11] Frank R. Hampel. The influence curve and its role in robust estimation. *Journal of the American Statistical Association*, 69(346):383–393, 1974.
- [12] Richard M. Karp. Reducibility among combinatorial problems. In Raymond E. Miller and James W. Thatcher, editors, *Complexity of Computer Computations*, pages 85–103. Plenum Press, New York, 1972.
- [13] Nikolas Kuschnig, Gregor Zens, and Jesús Crespo Cuaresma. Hidden in plain sight: influential sets in linear models. CESifo Working Paper 8981, CESifo, Munich, 2021.
- [14] Nicolai Meinshausen and Peter Bühlmann. Stability selection. *Journal of the Royal Statistical Society: Series B*, 72(4):417–473, 2010.
- [15] Ankur Moitra and Dhruv Rohatgi. Provably auditing ordinary least squares in low dimensions, 2022. arXiv:2205.14284; International Conference on Learning Representations 2023.

- [16] Michael R. Osborne, Brett Presnell, and Berwin A. Turlach. On the LASSO and its dual. *Journal of Computational and Graphical Statistics*, 9(2):319–337, 2000.
- [17] Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. Scikit-learn: machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [18] Jack Sherman and Winifred J. Morrison. Adjustment of an inverse matrix corresponding to a change in one element of a given matrix. *The Annals of Mathematical Statistics*, 21(1):124–127, 1950.
- [19] Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B*, 58(1):267–288, 1996.
- [20] Ryan J. Tibshirani. The lasso problem and uniqueness. *Electronic Journal of Statistics*, 7:1456–1490, 2013.
- [21] Martin J. Wainwright. Sharp thresholds for high-dimensional and noisy sparsity recovery using ℓ_1 -constrained quadratic programming (Lasso). *IEEE Transactions on Information Theory*, 55(5):2183–2202, 2009.
- [22] Max A. Woodbury. Inverting modified matrices. Memorandum Report 42, Statistical Research Group, Princeton University, Princeton, 1950.
- [23] Peng Zhao and Bin Yu. On model selection consistency of lasso. *Journal of Machine Learning Research*, 7:2541–2563, 2006.